

Lecture notes

Advanced Programming

WS 2018/19

Prof. Dr. Michael Hanus
Priv.Do. Dr. Frank Huch

Programming Languages and Compiler Construction
Department of Computer Science
Kiel University

November 6, 2018

Pre-Preface

These lecture notes are a translation of the german lecture notes associated with the “Advanced Programming” lectures held by Prof. Michael Hanus. Moreover, these notes do not translate the entire lecture by Prof. Hanus but only the chapters relevant according to the newly established pre-masters conditions for international students in Computer Science at the University of Kiel.

Preface

This lecture introduces advanced programming concepts which go beyond the first semesters of programming. The following is presented on the basis of various programming languages that represent the most important programming paradigms. Modern functional programming techniques are developed using the language Haskell. Logical and constraint-oriented programming is shown in the language Prolog. Concepts for concurrent and distributed programming are presented using the language Java.

This script is a revised version of a transcript, originally written and typeset in L^AT_EX by Nick Prühs in SS 2009. My thanks go to Nick Prühs for the first L^AT_EX template as well as Björn Peemöller and Lars Noelle for corrections.

One more important note: This script is only intended to give an overview of what is covered in the lecture. It does not replace attending the lecture, which is important for understanding the concepts and techniques of advanced programming. For in-depth self-study, a look at the textbooks and references given in the lecture is recommended.

Kiel, February 2017

Michael Hanus

PS: Anyone who did not find any errors while reading the notes probably has not paid enough attention. I am thankful for all hints on errors I receive, both in person as well as in writing or by e-mail.

Contents

1	Java Generics	1
1.1	Interaction with inheritance	3
1.2	Wildcards	4
1.3	Type inference	6
2	Concurrent programming in Java	7
2.1	Synchronization	8
2.1.1	Interprocess communication and synchronization	8
2.1.2	Synchronization with semaphores	8
2.1.3	Dining philosophers	10
2.2	Threads in Java	12
2.2.1	The class <code>Thread</code>	12
2.2.2	The interface <code>Runnable</code>	13
2.2.3	Properties of Thread objects	13
2.2.4	Synchronization of threads	14
2.2.5	The example class <code>Account</code>	14
2.2.6	Closer look at <code>synchronized</code>	15
2.2.7	Differentiating synchronization in the context of OO	17
2.2.8	Communication between threads	18
2.2.9	Case study: single-element buffer	21
2.2.10	Exiting and interrupting threads	23
2.3	Distributed programming in Java	24
2.3.1	Serialization/Deserialization of data	24
2.3.2	Remote Method Invocation (RMI)	25
2.3.3	RMI registry	27
3	Functional Programming	30
3.1	Expressions and functions	31
3.1.1	Evaluation	33
3.1.2	Local definitions	35
3.2	Data types	37
3.2.1	Basic data types	37
3.2.2	Type annotations	38
3.2.3	Algebraic data structures	39
3.3	Polymorphism	41
3.4	Pattern Matching	45
3.4.1	Structure of patterns	46

Contents

3.4.2	Case expressions	46
3.4.3	Guards	47
3.5	Higher Order Functions	48
3.5.1	Example: derivative function	48
3.5.2	Anonymous functions (lambda abstraction)	48
3.5.3	Generic Programming	51
3.5.4	Control structures	53
3.5.5	Functions as data	54
3.5.6	Useful Higher Order Functions	55
3.5.7	Higher Order Functions in Imperative Languages	56
3.6	Type Classes and Overloading	60
3.6.1	Predefined Functions of a Class	61
3.6.2	Predefined Classes	61
3.6.3	The class <code>Read</code>	62
3.7	Evaluation Strategies and Lazy Evaluation	64
Bibliography		77
List of Figures		79

1 Java Generics

Since version 5.0 (released in 2004) Java supports *generic programming*: Classes and methods can be parameterized with types. This opens up similar possibilities as templates in C++. As an easy example we consider a very simple container class which can store either a value or no value. In Java, this could be defined as follows.

```
public class Optional {  
  
    private Object value;  
    private boolean present;  
  
    public Optional() {  
        present = false;  
    }  
  
    public Optional(Object v) {  
        value = v;  
        present = true;  
    }  
  
    public boolean isPresent() {  
        return present;  
    }  
  
    public Object get() {  
        if (present) {  
            return value;  
        }  
        throw new NoSuchElementException();  
    }  
}
```

We use the fact that `NoSuchElementException` is derived from `RuntimeException`, that is, it is a so-called *unchecked exception* and does not need to be declared.

The class could be used as follows.

```
Optional opt = new Optional(new Integer(42));  
  
if (opt.isPresent()) {  
    Integer n = (Integer) opt.get();  
}
```

```

        System.out.println(n);
    }

```

Each time the stored object of the container class is accessed, an explicit type cast is done. There is *no type security*: In case of a false cast a `ClassCastException` is thrown at runtime.

Keep in mind that the definition of the class `Optional` is independent of the type of the stored value; the type of `value` is `Object`. This allows using arbitrary types that are represented in an abstract way in the definition of the class: Passing a type as a parameter is called *parametric polymorphism*.

Type parameters are marked by angle brackets and are used in place of `Object`. The adapted class definition is.

```

public class Optional<T> {

    private T value;
    private boolean present;

    public Optional() {
        present = false;
    }

    public Optional(T v) {
        value = v;
        present = true;
    }

    public boolean isPresent() {
        return present;
    }

    public T get() {
        if (present) {
            return value;
        }
        throw new NoSuchElementException();
    }
}

```

The class is now used like this.

```

Optional<Integer> opt = new Optional<Integer>(new Integer(42));

if (opt.isPresent()) {
    Integer n = opt.get();
    System.out.println(n);
}

```

```
}
```

Explicit type casts are not necessary and an expression like

```
opt = new Optional<Integer>(opt);
```

returns a type error at compile time.

Naturally, multiple type parameters are also possible.

```
public class Pair<A, B> {

    private A first;
    private B second;

    public Pair(A first, B second) {
        this.first = first;
        this.second = second;
    }

    public A first() {
        return first;
    }

    public B second() {
        return second;
    }
}
```

1.1 Interaction with inheritance

It is also possible to restrict type parameters so that only classes can be used that provide certain methods. As an example we will extend the class `Optional` so that it implements the interface `Comparable`.

```
public class Optional<T extends Comparable<T>>
    implements Comparable<Optional<T>> {
    ...
    @Override
    public int compareTo(Optional<T> o) {
        if (present) {
            return o.isPresent() ? value.compareTo(o.get()) : 1;
        } else {
            return o.isPresent() ? -1 : 0;
        }
    }
}
```

```
}
}
```

For interfaces as well as inheritance the keyword `extends` is used. It is also possible to list multiple restrictions of type variables. For three type parameters (`T`, `S` and `U`) this would look like the following.

```
<T extends A<T>, S, U extends B<T,S>>
```

In this example, `T` has to provide the methods of `A<T>` and `U` has to provide the methods of `B<T,S>`. `A` and `B` need to be classes or interfaces, not type variables.

1.2 Wildcards

The class `Integer` is a sub class of the class `Number`. Therefore, the following code should be valid.

```
Optional<Integer> oi = new Optional<Integer>(new Integer(42));
Optional<Number> on = oi;
```

However, the type system does not allow this because `Optional<Integer>` is **no** subtype of `Optional<Number>`! The problem becomes more obvious when we add a setter method to the class `Optional`.

```
public void set(T v) {
    present = true;
    value = v;
}
```

In consequence, this would allow the following.

```
on.set(new Float(42));
Integer i = oi.get();
```

Since `oi` and `on` are names for the same objects, executing this code would lead to a type error. For this reason the type system does not allow the above code.

Nevertheless, sometimes a supertype of polymorphic classes, that is, a type that includes all other types, is needed. An `Optional` of unknown type is described like this.

```
Optional<?> ox = oi;
```

1 Java Generics

The symbol “?” is known as a *wildcard* type. `Optional<?>` represents all other instances of `Optional`, for example `Optional<Integer>`, `Optional<String>` and `Optional<Optional<Object>>`. One downside of this approach is that only methods, which fit every type, work.

```
Object o = ox.get();
```

We can use the method `set` with `ox` but we need a value that belongs to every type: `null`.

```
ox.set(null);
```

This is the only possible `set`-call.

Unfortunately, the wildcard type “?” is often not sufficient. For example, when multiple subtypes like GUI elements are stored in a collection. This can be solved by using *bounded wildcards*.

`<? extends A>` includes *all* subtypes of `A` (*covariance*)

`<? super A>` includes *all* supertypes of `A` (*contra variance*)

Keep in mind that a type is sub- and supertype of itself, that is, for all types `T` the properties `<T extends T>` and `<T super T>` hold.

In the following, a few examples of using bounded wildcards are shown.

Expression	valid?	Reasoning
<hr/>		
<code>Optional<? extends Number> on = oi;</code>		
<code>Integer i = on.get();</code>	no!	? could also be <code>Float</code>
<code>Number n = on.get();</code>	yes	
<code>on.set(n);</code>	no!	? could be more specific than <code>Number</code>
<code>on.set(new Integer(42));</code>	no!	? could also be <code>Float</code>
<code>on.set(null);</code>	yes	
<hr/>		
<code>Optional<? super Integer> ox = oi;</code>		
<code>Integer i = ox.get();</code>	no!	? could also be <code>Number</code> or <code>Object</code>
<code>Number n = ox.get();</code>	no!	? could also be <code>Object</code>
<code>Object o = ox.get();</code>	yes	
<code>ox.set(o);</code>	no!	? could also be <code>Number</code>
<code>ox.set(n);</code>	no!	<code>n</code> could also have the type <code>Float</code>
<code>ox.set(i);</code>	yes	

Figure 1.1: Expressions with wildcards

This means that we can extract objects of type `A` from an `Optional<? extends A>` but only insert `null`, while we can insert objects of the type `A` into an `Optional<? super A>`

but only extract objects of type `Object`.

1.3 Type inference

When initializing a variable with a generic type parameter, the parameter is stated twice.

```
Optional<Integer> o = new Optional<Integer>(new Integer(42));  
List<Optional<Integer>> l = new ArrayList<Optional<Integer>>(o);
```

Since version 7, the Java compiler is able to infer the generic type when initializing an object with `new` in most cases. Therefore, the type only needs to be stated in the variable declaration. When calling the `new` operator, the type is calculated by the diamond operator `<>`.

```
Optional<Integer> mv = new Optional<>(new Integer(42));  
List<Optional<Integer>> l = new ArrayList<>(mv);
```

Only the whole type within the angle brackets can be omitted, types like `<Optional<>>` are not allowed. If the compiler is not able to infer the type, it still needs to be stated manually.

Besides the shorter code, the diamond operator also enables creating generic singleton objects, which was not possible before.

```
Optional<?> empty = new Optional<>();
```

This is especially useful when only one instance of immutable objects should be kept in memory.

2 Concurrent programming in Java

An application in computer science is described as *concurrent* if it does not have a strictly sequential flow, but when the application has several activities which run concurrently, that is, (almost) in parallel. These activities are often referred to as tasks, threads, or processes. These tasks are sequentially running programs themselves. In a concurrent program, there are therefore not only one program counter which determines the next instruction to be processed but many program counters. We will see that the development of concurrent software is associated with many pitfalls. In this chapter we will discuss how to develop concurrent software that is as reliable as possible. Building on concurrent concepts, we will later also consider distributed software. We will use Java as programming language, but the concepts can also be applied to other languages.

Why do we need concurrent programming? We would often like one application to take over several tasks. At the same time the *reactivity* of the application should be preserved. Examples of such applications are listed in the following.

- GUIs
- operating system routines
- distributed applications (web servers , chat, ...)

Solution We achieve this by means of *Concurrency*). By the use of threads or processes, individual tasks of an application can be programmed and executed independently of other tasks.

Parallelism The parallel execution of several processes intends to achieve faster execution (high-performance computing).

Distributed system Several components in a network work on a problem together. Usually there is a distributed task, sometimes distributed systems are used for parallelization.

When we speak of *multitasking*, we mean that the processor time is distributed by a scheduler among concurrent threads or processes. We distinguish between two types of multitasking.

Cooperative multitasking A thread continues to calculate until it returns control (for example with `yield()`) or until it waits for messages (`suspend()`). In Java we find this in so-called *green threads*.

Preemptive multitasking The scheduler can take control from tasks. Here we often enjoy more programming comfort because we do not need to think about where we should give up control.

2.1 Synchronization

2.1.1 Interprocess communication and synchronization

In addition to the generation of threads or processes, communication between them is also important. This is usually done via shared memory or variables. We consider the following example in pseudo code.

```
int i = 0;
par
  { i = i + 1; }
  { i = i * 2; }
end par;
print(i);
```

Concurrency makes programs non-deterministic, that is, different results can be obtained depending on the scheduling. Thus, the above program can generate outputs 1 or 2, depending on how the scheduler executes the two concurrent processes. If the result of a program run depends on the order of the scheduling, this is called a *race condition*.

In addition to the two results 1 and 2, another result, namely 0, is also possible. This is due to the fact that it has not yet been clearly specified which actions are really executed atomically. By translating the program into byte or machine code, the following instructions can result.

1. `i = i + 1;` \rightarrow `LOAD i; INC; STORE i;`
2. `i = i * 2;` \rightarrow `LOAD i; SHIFTL; STORE i;`

Then the following sequence leads to output 0.

- ```
(2) LOAD i;

(1) LOAD i;
(1) INC;
(1) STORE i;

(2) SHIFTL;
(2) STORE i;
```

So we need synchronization to ensure the atomic execution of certain sections of code that work concurrently on the same resources. We call such code sections *critical sections*.

### 2.1.2 Synchronization with semaphores

A well-known concept for synchronizing concurrent threads or processes goes back to Dijkstra from 1968. Dijkstra developed an abstract data type with the aim of synchronizing the atomic (uninterrupted) execution of certain program sections. These *semaphores* provide two atomic operations.

```

p(s) {
 if s >= 1
 then s = s - 1;
 else add executing thread to queue for s and suspend it;
}

```

```

v(s) {
 if waiting list for s not empty
 then wake first process in queue
 else s = s + 1;
}

```

$p(s)$  stands for pass or *passeer*,  $v(s)$  stands for leave or *verlaat*. Now we can prevent the output 0 of our above program as follows.

```

int i = 0;
Semaphore s = 1;
par
 { p(s); i = i + 1; v(s); }
 { p(s); i = i * 2; v(s); }
end par;

```

The initial value of the semaphore determines the maximum number of processes in the critical area. Usually we find the value 1 here. We also call such semaphores *binary semaphore*.

Another application of semaphores is the *producer-consumer problem*:  $n$  producers produce goods that are consumed by  $m$  consumers. One simple solution to this problem uses an unrestricted buffer.

---

```

Buffer buffer = ...
Semaphore num = 0;

```

---

The producer's code looks like this.

---

```

while (true) {
 newproduct = produce();
 push(newproduct, buffer);
 v(num);
}

```

---

The consumer's code looks like this.

---

```

while (true) {
 p(num);
 prod = pull(buffer);
 consume(prod);
}

```

---

```
}
```

---

What is still missing is the synchronization to `buffer`. The synchronization can be realized by adding another semaphore.

---

```
Buffer buffer = ...
Semaphore num = 0;
Semaphore bufferAccess = 1;
```

---

Below is the adapted code of the producer.

---

```
while (true) {
 newproduct = produce();
 p(bufferAccess);
 push(newproduct, buffer);
 v(bufferAccess);
 v(num);
}
```

---

Below is the adapted code of the consumer.

---

```
while (true) {
 p(num);
 p(bufferAccess);
 prod = pull(buffer);
 v(bufferAccess);
 consume(prod);
}
```

---

However, the use of semaphores also has some disadvantages. The code with semaphores quickly looks unstructured and confusing. In addition, we cannot use semaphores compositionally: the simple code `p(s); p(s);` can already generate a *deadlock* on a binary semaphore `s`.

The concept of *monitors* which may be familiar from lectures like “Operating and Communication Systems” offers an improvement here. In fact, Java uses a mechanism similar to these monitors for synchronization.

### 2.1.3 Dining philosophers

The *dining philosophers* problem with  $n$  philosophers can be modelled as follows using semaphores.

---

```
Semaphore stick1 = 1;
Semaphore stick2 = 1;
```

```
Semaphore stick3 = 1;
Semaphore stick4 = 1;
Semaphore stick5 = 1;

par { phil(stick1, stick2); }
 { phil(stick2, stick3); }
 { phil(stick3, stick4); }
 { phil(stick4, stick5); }
 { phil(stick5, stick1); }
end par;
```

---

The code for philosopher *i* looks like the following.

---

```
public phil(stickl,stickr) {
 while (true) {
 think();

 p(stickl);
 p(stickr);

 eat();

 v(stickl);
 v(stickr);
 }
}
```

---

However, a deadlock can occur if all philosophers take their left stick at the same time. We can avoid this deadlock by putting it back.

---

```
while (true) {
 think();

 p(stickl);

 if (lookup(stickr) == 0) { # look up integer value of stick
 v(stickl);
 } else {
 p(stickr);

 eat();

 v(stickl);
 v(stickr);
 }
}
```

---

Here `lookup(s)` denotes a lookup function of the abstract data type semaphore that returns the integer value of a semaphore `s`.

The program now has a livelock, that is, individual philosophers can starve to death. We do not want to discuss this problem further here.

## 2.2 Threads in Java

### 2.2.1 The class Thread

The API of Java offers a class `Thread` in the package `java.lang`. Own threads can be derived from this. The code to be executed in parallel is written to the `run()` method. After we have created a new thread with the help of its constructor, we can start it for concurrent execution with the method `start()`.

We consider the following simple thread as an example.

---

```
public class ConcurrentPrint extends Thread {
 private String s;

 public ConcurrentPrint(String s) {
 this.s = s;
 }

 public void run() {
 while (true) {
 System.out.print(s + " ");
 }
 }

 public static void main(String[] args) {
 new ConcurrentPrint("a").start();
 new ConcurrentPrint("b").start();
 }
}
```

---

The above program flow can lead to many possible outputs:

```
a a b b a a b b ...
a a a b b ...
a b a a a b a a b b ...
a a a a a a a a a a ...
```

The latter is guaranteed if cooperative scheduling is used.



### 2.2.2 The interface Runnable

Java does not offer inheriting from multiple classes. Therefore, an extension of the class `Thread` is often unfavorable. An alternative is the Interface `Runnable`:

---

```
public class ConcurrentPrint implements Runnable {
 private String s;

 public ConcurrentPrint(String s) {
 this.s = s;
 }

 public void run() {
 while (true) {
 System.out.print(s + " ");
 }
 }

 public static void main(String[] args) {
 Runnable aThread = new ConcurrentPrint("a");
 Runnable bThread = new ConcurrentPrint("b");

 new Thread(aThread).start();
 new Thread(bThread).start();
 }
}
```

---

Note that within the above implementation of `ConcurrentPrint`, `this` no longer returns an object of type `Thread`. The current thread object can instead be reached via the static method `Thread.currentThread()`.

### 2.2.3 Properties of Thread objects

Every thread object in Java has a number of properties.

**Name** Each thread has a name, such as "main-Thread", "Thread-0" or "Thread-1". Access to the name of a thread is defined using the methods `getName()` and `setName(String)`. You can use these for example as names for debugging.

**State** Each thread is always in a certain state. An overview of these states and the state transitions is shown in figure 2.1. A thread object remains in the *terminated* state until all references to it have been discarded.

**Demon** A thread can be created as a background thread by calling `setDaemon(true)` before calling the `start()` method. The JVM terminates when only daemon threads are running. Examples of such threads are AWT threads or the garbage collector.

**Priority** Each thread in Java has a specific priority that is platform-specific.

**Thread group** Threads can also be divided into groups for simultaneous management.




Figure 2.1: States of threads

**Method** `sleep(long)` Lets the thread sleep for the specified time. A call to this method can throw a `InterruptedException` which must be caught.

### 2.2.4 Synchronization of threads

Java offers a monitor-like concept for thread synchronization that allows to set and release locks on objects.

The methods of a thread in Java can be declared as `synchronized`. In all synchronized methods of an object there may be a maximum of one thread at a time. This also includes calculations that are called in a synchronized method and unsynchronized methods of the same object. A synchronized method is not left by calling `sleep(long)` or `yield()`.

Furthermore, each object has its own *lock*. When attempting to execute a method declared as `synchronized`, we distinguish three cases.

1. If the lock is released, the thread takes it.
2. If the thread already has the lock, it continues.
3. Otherwise, the thread is suspended.

The lock is released again if the method is exited in which it was acquired. Compared to semaphores, the Java approach seems more structured, you can't forget to unlock. Yet it is less flexible.

### 2.2.5 The example class `Account`

A simple example is intended to illustrate the use of `synchronized` methods. We are looking at an implementation of a class for a bank account.

---

```
public class Account {
 private int balance;

 public Account(int initialDeposit) {
 balance = initialDeposit;
 }

 public synchronized int getBalance() {
 return balance;
 }

 public synchronized void deposit(int amount) {
 balance += amount;
 }
}
```

---

We now want to use the class as follows.

---

```
Account a = new Account(300);
...
a.deposit(100); // concurrent, first thread
...
a.deposit(100); // concurrent, second thread
...
System.out.println(a.getBalance());
```

---

The calls of the method `deposit(int)` should be made concurrently from different threads. Without the keyword `synchronized`, the output would be 500. The output 400 would also be possible (see section 2.1.1), Interprocess communication and synchronization). With `synchronized`, an output of 500 is guaranteed.

### 2.2.6 Closer look at synchronized

Inherited synchronized methods do not necessarily have to be synchronized again. If you overwrite such methods, you can omit the keyword `synchronized`. This is called *refined implementation*. The method of the superclass remains `synchronized`. On the other hand, unsynchronized methods can also be overwritten by synchronized ones.

Class methods declared as synchronized (`static synchronized`) have no interaction with synchronized object methods. The class therefore has its own lock.

You can also synchronize individual statements.

---

```
synchronized (expr) block
```

---

The expression `expr` must evaluate an object, whose lock is then used for synchronization. Strictly speaking, synchronized methods are only syntactic sugar: The method declaration

---

```
synchronized A m(args) block
```

---

is expanded to

---

```
A m(args) {
 synchronized (this) block
}
```

---

Synchronizing individual statements is useful to reduce the amount of code that needs to be synchronized or sequentialized.

---

```
private double state;

public void calc() {
 double res;
 // do some really expensive computation
 ...
 // save the result to an instance variable
 synchronized (this) {
 state = res;
 }
}
```

---

Synchronizing individual statements is also useful to synchronize on other objects. We consider a simple implementation of a synchronized collection as an example.

---

```
class Store {
 public synchronized boolean hasSpace() {
 ...
 }

 public synchronized void insert(int i)
 throws NoSpaceAvailableException {
 ...
 }
}
```

---

We now want to use this collection as follows.

---

```
Store s = new Store();
```

```
...
if (s.hasSpace()) {
 s.insert(42);
}
```

---

But this leads to problems, because we cannot exclude that a re-schedule happens between the calls of `hasSpace()` and `insert(int)`. Since defining special methods for such cases often turns out to be impracticable, we better use the collection like this.

---

```
synchronized(s) {
 if (s.hasSpace()) {
 s.insert(42);
 }
}
```

---

### 2.2.7 Differentiating synchronization in the context of OO

We call synchronized methods and synchronized instructions in object methods *server-side synchronisation*. Synchronization of the calls of an object is called *client-side synchronisation*.

For efficiency reasons, Java API objects, especially collections, are no longer synchronized. For Collections, however, there exist synchronized versions via wrappers such as `synchronizedCollection`, `synchronizedSet`, `synchronizedSortedSet`, `synchronizedList`, `synchronizedMap` or `synchronizedSortedMap`.

Safely copying a list into an array is possible in two different ways now. First, we create an instance of a synchronized list.

---

```
List<Integer> unsyncList = new List<Integer>();
... // fill the list
List<Integer> list = Collections.synchronizedList(unsyncList);
```

---

Now we copy this list to an array with either a single line

---

```
Integer[] a = list.toArray(new Integer[0]);
```

---

or with

---

```
Integer[] b;

synchronized (list) {
 b = new Integer[list.size()];
 list.toArray(b);
}
```

---

}

With the second, two-line variant, synchronization to the list is indispensable: We access the collection in both lines and cannot guarantee that another thread will not change the collection in the meantime. This is a classic example of client-side synchronization.

### 2.2.8 Communication between threads

Threads communicate with each other via shared objects. How to find we find out when a variable contains a value? For this there are several possible solutions.

The first option is to display the modification of a component of the object, for example, by setting a flag (`boolean`). However, this has the disadvantage that checking the flag leads to busy waiting. Busy waiting means that a thread is waiting for the occurrence of an event, thereby continuing to calculate and thus consuming resources like processor time. Therefore, a thread is suspended using the method `wait()` of the object and woken up later by using `notify()` or `notifyAll()`.

---

```
public class C {
 private int state = 0;

 public synchronized void printNewState()
 throws InterruptedException {

 wait();
 System.out.println(state);
 }

 public synchronized void setValue(int v) {
 state = v;
 notify();
 System.out.println("value set");
 }
}
```

---

Two thread call the methods `printNewState()` and `setValue(42)` concurrently. Now the *only* possible output is

---

```
value set
42
```

---

If the call of `wait()` only comes after the method `setValue(int)` has already been left by the first thread, this leads to the output `value set`.

The methods `wait()`, `notify()` and `notifyAll()` may only be used within `synchronized` methods or blocks and are methods of the locked object with the following semantics.

`wait()` puts the executing thread to sleep and releases the lock of the object again. `notify()` awakens *one* sleeping thread of the object and continues with its own calculation. The awakened thread now applies for the lock. If no thread sleeps, the `notify()` is lost.

`notifyAll()` does the same as `notify()`, only for all threads that were laid to sleep with `wait()` for this object.

Please note that these three methods may only be called on objects whose lock has been received before. The call must therefore be made in a `synchronized` method or in a `synchronized` block, otherwise a `IllegalMonitorStateException` is thrown at runtime.

We would like to write a program that outputs all changes of the state.

---

```
...
private boolean modified = false; // signals state changes
...
public synchronized void printNewState()
 throws InterruptedException {
 while (true) {
 if (!modified) {
 wait();
 }

 System.out.println(state);
 modified = false;
 }
}

public synchronized void setValue(int v) {
 state = v;
 notify();
 modified = true;
 System.out.println("value set");
}
```

---

One thread now executes `printNewState()`, other threads change the state using `setValue(int)`. This leads to a problem: With several setting threads, the output of individual intermediate states can be lost. So `setValue(int)` also has to wait and wake up if necessary.

---

```
public synchronized void printNewState()
 throws InterruptedException {
 while (true) {
 if (!modified) {
 wait();
 }

 System.out.println(state);
 }
}
```

---

```

 modified = false;
 notify();
 }
}

public synchronized void setValue(int v)
 throws InterruptedException {
 if (modified) {
 wait();
 }

 state = v;
 notify();
 modified = true;
 System.out.println("value set");
}

```

---

But now it is not guaranteed that the call of `notify()` in the method `setValue(int)` wakes up the `printNewState` thread! In Java we solve this problem with `notifyAll()` and accept a little busy waiting.

---

```

public synchronized void printNewState()
 throws InterruptedException {
 while (true) {
 while (!modified) {
 wait();
 }

 System.out.println(state);
 modified = false;
 notify();
 }
}

public synchronized void setValue(int v)
 throws InterruptedException {
 while (modified) {
 wait();
 }

 state = v;
 notifyAll();
 modified = true;
 System.out.println("value set");
}

```

---

The `wait()` method is also overloaded several times in Java:



`wait(long)` interrupts execution for the specified number of milliseconds.

`wait(long, int)` interrupts the execution for the specified number of milli- and nanoseconds.

Note: It is strongly discouraged to base the correctness of the program on these overloads! The calls `wait(0)`, `wait(0, 0)` and `wait()` all cause the thread to wait until it wakes up again.

### 2.2.9 Case study: single-element buffer

A single-element buffer is convenient for communication between threads. Since the buffer is single-element, it can only be empty or full. A value can be written into an empty buffer via a method `put` from a full buffer. The value can be removed using `take`. `take` suspends on an empty buffer, `put` suspends on a full buffer.

---

```
public class Buffer1<T> {
 private T content;
 private boolean empty;

 public Buffer1() {
 empty = true;
 }

 public Buffer1(T content) {
 this.content = content;
 empty = false;
 }

 public synchronized T take() throws InterruptedException {
 while (empty) {
 wait();
 }

 empty = true;
 notifyAll();

 return content;
 }

 public synchronized void put(T o) throws InterruptedException {
 while (!empty) {
 wait();
 }

 empty = false;
 }
}
```

```

 notifyAll();
 content = o;
}

public synchronized boolean isEmpty() {
 return empty;
}
}

```

---

What is unfortunate about the above solution is that too many threads are awakened, that is, `notifyAll()` always awakens all reading threads as well as all writing threads, most of which are immediately put to sleep again. Can we wake up threads in a targeted way? Yes! We use special objects to synchronize the `take` and `put` threads.

---

```

public class Buffer1<T> {
 private T content;
 private boolean empty;
 private Object r = new Object();
 private Object w = new Object();

 public Buffer1() {
 empty = true;
 }

 public Buffer1(T content) {
 this.content = content;
 empty = false;
 }

 public T take() throws InterruptedException {
 synchronized (r) {
 while (empty) {
 r.wait();
 }

 synchronized (w) {
 empty = true;
 w.notify();

 return content;
 }
 }
 }

 public void put(T o) throws InterruptedException {
 synchronized (w) {

```

```

 while (!empty) {
 w.wait();
 }

 synchronized (r) {
 empty = false;
 r.notify();
 content = o;
 }
}

public boolean isEmpty() {
 return empty;
}
}

```

---

Here the `while` is very important! Another thread entering the method from the outside could overtake a waiting (and just awakened) thread!

### 2.2.10 Exiting and interrupting threads

Java offers several ways to terminate threads:

1. terminating the `run()` method
2. aborting the `run()` method
3. calling the `destroy()` method (deprecated, partly no longer implemented)
4. demon thread and program end

With 1 and 2 all locks are released. With 3, locks are not released which makes this method uncontrollable. For this reason, this method should not be used. With 4, the locks do not matter.

Java also provides a way to interrupt threads via *interrupts*. Each thread has a flag indicating interrupts.

The `Thread` method `interrupt()` sends an interrupt to a thread, the flag is set. If the thread is sleeping due to a call to `sleep()` or `wait()`, it is awakened and a `InterruptedException` is thrown.

---

```

synchronized (o) {
 ...
 try {
 ...
 o.wait();
 ...
 } catch (InterruptedException e) {
 ...
 }
}

```

```
}
}
```

---

In an interrupt after calling `wait()`, the `catch` block is not entered until the thread has recovered the lock on the `o` object of the surrounding `synchronized` block!

In contrast, in the suspension with `synchronized` the thread is not awakened, but only the flag is set.

The method `public boolean isInterrupted()` tests if a thread has received interrupts. `public static boolean interrupted()` tests the current thread for an interrupt and clears the interrupted flag.

Handling interrupts in a `synchronized` method is possible as follows:

---

```
synchronized void m(...) {
 ...
 if (Thread.currentThread().isInterrupted()) {
 throw new InterruptedException();
 }
}
```

---

If a `InterruptedException` is caught, the flag is also cleared. Then the flag needs to be set again!

## 2.3 Distributed programming in Java

As an abstraction of network communication, Java offers the *Remote Method Invocation (RMI)*. This allows remote objects to be used on other computers as if they were local objects. In order for the data to be sent over a network, it must be converted (arguments and results of method calls) into byte sequences, usually referred to as serialization.

### 2.3.1 Serialization/Deserialization of data

The serialization of an object `o` returns a byte sequence, the deserialization of the byte sequence returns a new object `o1`. Both objects should be the same with regard to their behavior, but have different object identities, i.e. `o1` is a copy of `o`.

The (de-)serialization takes place recursively. Thus, contained objects must also be (de)serialized. To specify that an object can be serialized, Java offers the interface `Serializable`.

---

```
public class C implements java.io.Serializable { ... }
```

---

This interface does not contain any methods and is therefore only a "marker interface" that specifies that objects of this class can be (de-)serialized.

During serialization, certain time- or security-critical parts of an object can be hidden using the keyword `transient` :

---

```
protected transient String password;
```

---

Transient values should be set explicitly after deserialization (for example a timer) or should not be used (for example passwords).

For serializing you use the class `ObjectOutputStream` , for deserializing the class `ObjectInputStream` . Their constructors are a `OutputStream` and `InputStream` respectively. Objects can be written

---

```
public void writeObject(Object o)
```

---

using and read by means of

---

```
public final Object readObject()
```

---

and a cast of appropriate type.

This could be used to "manually" transfer objects from one machine to another (for example by using socket connections) and then work with them on the other machine. You can avoid this manual work, however, if you only want to use certain functionalities of an object somewhere else. This is made possible in Java by Remote Method Invocation.

### 2.3.2 Remote Method Invocation (RMI)

In OO-programming we have a client-server view of objects. When a method is called, the calling object is seen as the client and the called object as the server.

In a distributed context, messages become real messages on the Internet (transmitted via TCP). Processes that communicate with each other are

- servers that provide information.
- clients that request information.

Any communication pattern (for example peer-to-peer) can be represented. The idea of RMI goes back to *Remote Procedure Call (RPC)*, which was developed for C.

First, an RMI client needs a reference to the remote object. The RMI registration serves this purpose. It requests the reference using a URL:

```
rmi://hostname:port/servicename
```

`hostname` can be a computer name or an IP address, `servicename` is a string describing an object. The default port of RMI is 1099.

There is also a second way to access a remote object in a program: as the argument of a method call. Usually you can use the above registration only for the "first contact", then (remote) objects are exchanged and used transparently (like local objects).




Figure 2.2: Remote Method Invocation in Java

Such objects can be distributed on any computer. Method calls for remote objects are implemented by network communication.

The interface for RMI is divided into stub and skeleton. Since Java 5 these are no longer visible, but are generated implicitly at runtime.

The network communication takes place via TCP/IP, but this is also not visible for the application programmer.

The parameters and the return value of a remotely called method must of course be converted into bytes and be transferred. The following rules apply.

- For `Remote` objects only a reference of the object is transferred.
- `Serializable` objects are converted to `byte[]`. On the other side a copy of the object will be created.
- Primitive values are copied.

To use objects from remote nodes, an RMI Service Interface must first be defined for this object. This interface describes the methods that can be called from another computer. Consider a simple “Flip” server as an example, that is, an object containing a Boolean state which can be changed by a `flip` message. The RMI service interface can then be defined as follows.

---

```
import java.rmi.*;

public interface FlipServer extends Remote {
 public void flip() throws RemoteException;
 public boolean getState() throws RemoteException;
}
```

---

An implementation of the RMI interface on the server side could look like the following.

---

```
import java.rmi.server.*;
import java.rmi.*;

public class FlipServerImpl extends UnicastRemoteObject
 implements FlipServer {

 private boolean state;

 public FlipServerImpl() throws RemoteException {
 state = false;
 }

 public void flip() {
 state = !state;
 }

 public boolean getState() {
 return state;
 }
}
```

---

In older Java versions it was necessary to generate the stub and skeleton classes with the RMI compiler `rmic` - this is no longer the case. The class `UnicastRemoteObject` represents the simplest way of realizing remote objects. All necessary translation steps are done automatically. In some cases, it may be necessary to intervene in the (de-)serialization process yourself to define properties of the `RemoteObjects`. For this purpose, other interfaces exist.

### 2.3.3 RMI registry

To access a server object from another host, Java provides an RMI registry. To access a remote object, an RMI registry server must be initialized on the host on which the server object is executed. This RMI registry server can be started explicitly using the command `rmiregistry` (for example, in a UNIX shell). Its use for object registration is shown in the following example.

---

```
import java.rmi.*;
import java.net.*;

public class Server {
 public static void main(String[] args) {
 try {
 String host = args.length >= 1 ? args[0] : "localhost";
 String uri = "rmi://" + host + "/FlipServer";

 FlipServerImpl server = new FlipServerImpl();
 Naming.rebind(uri, server);
 } catch (MalformedURLException|RemoteException e) { ... }
 }
}
```

---

rebind registers the name of the server object. On the client side, the RMI registry must be contacted to get a reference to the remote object so that its flip method can be called.

---

```
import java.rmi.*;
import java.net.*;

public class Client {
 public static void main(String[] args) {
 try {
 String host = args.length >= 1 ? args[0] : "localhost";
 String uri = "rmi://" + host + "/FlipServer";

 FlipServer s = (FlipServer) Naming.lookup(uri);
 s.flip();
 System.out.println("State: " + s.getState());
 } catch (MalformedURLException|NotBoundException|RemoteException e) {
 ...
 }
 }
}
```

---

However, remind that the RMI registry is only used for a "first contact" und usually only on server object is registered. After that, other remote objects are passed as prameters or return values and nodes in the distributed system get aware of all relevant objects of the whole system.

Thus, RMI represents a continuation of the sequential programming on distributed objects. This means that several distributed processes can access an object "simultaneously". So you also have to consider the problem of concurrent synchronization! However, synchronization can be more difficult here, as the following example shows.



As we have seen above, *client-side synchronization* is to guarantee the atomic execution of several methods which can already be synchronized. For this purpose we consider the simple Flip-server once again and a client that uses the `flip` method twice without interrupt calls.

---

```
...
FlipServer s = (FlipServer) Naming.lookup(uri);
synchronized(s) {
 System.out.println("State1: " + s.getState());
 s.flip();
 Thread.sleep(2000);
 s.flip();
 System.out.println("State2: " + s.getState());
}
...
```

---

Since the client synchronizes both calls on the flip server `s`, no one else should be able to change the state of the flip server during this time, so that the results of both outputs are always the same. This would also be the case in the purely concurrent context, but in the distributed context this no longer applies, since synchronization does not take place on the remote objects, but on the local stub objects!

Dynamic loading, security concepts or distributed memory cleanup (garbage collection) are other aspects of Java RMI that we do not consider here.

# 3 Functional Programming

One focus of this lecture is this chapter on functional programming. The practical relevance of functional programming is emphasized, for example, by Thomas Ball and Benjamin Zorn of Microsoft in the article [3], which bears the meaningful subtitle “Industry is ready and waiting for more graduates educated in the principles of programming languages”.

Second, would-be programmers (CS majors or non-majors) should be exposed as early as possible to functional programming languages to gain experience in the declarative programming paradigm. The value of functional/declarative language abstractions is clear: they allow programmers to do more with less and enable compilation to more efficient code across a wide range of runtime targets.

Another interesting recommendation of the authors is the following.

First, computer science majors, many of whom will be the designers and implementers of next-generation systems, should get a grounding in logic, its application in design formalisms, and experience the creation and debugging of formal specifications with automated tools. . .

Therefore, the study of mathematical foundations and logic is an important aspect of a computer science degree, but it is not part of this lecture.

Functional programming offers a number of advantages over classical imperative programming.

- High level of abstraction, no manipulation of memory cells
- No side effects, therefore easier code optimization and better comprehensibility
- Programming via properties, not via the execution sequence
- Implicit memory management
- Simpler proof of correctness and verification
- Compact source programs, therefore shorter development time, more readable programs, better maintainability
- Modular program structure, polymorphism, higher-order functions, reusability of code

For practical programming, knowledge of functional programming is important, as functional programming techniques and language constructs lead to better structured programs, as explained in the article [15]. Therefore, functional concepts can also be found in

many modern programming languages in a limited form. However, we will first introduce purely functional programming.

In functional programs, a *variable* represents an unknown value. A *program* is a set of function definitions. The memory is not explicitly usable, but is automatically managed and cleared. A *program flow* consists of the reduction of expressions. This goes back to the mathematical theory of the  $\lambda$  calculus from Church [5]. In the following we introduce the purely functional programming using the Haskell programming language [13].

## 3.1 Expressions and functions

In mathematics, a variable represents unknown (arbitrary) values, so that we often use expressions like

$$x^2 - 4x + 4 = 0 \Leftrightarrow x = 2$$

In imperative languages, on the other hand, we often see expressions such as the following.

$$x = x + 1 \quad \text{or} \quad x := x + 1$$

This represents a contradiction to mathematics. In functional programming, as in mathematics, variables are interpreted as unknown values (and not as names for memory cells)!

While functions in mathematics are used for calculation, we use procedures or functions in programming languages for structuring. There, however, is no real connection due to side effects. In functional programming languages, however, there are no side effects, so every function call with the same arguments produces the same result.

Functions can be defined in Haskell as follows.

---

```
f x1 ... xn = e
```

---

`f` is the function name, `x1` to `xn` are formal parameters or variables, and `e` is the body, an expression over `x1` to `xn`.

*Expressions* are constructed in Haskell by combining elements from the (incomplete) list below. gebildet werden:

1. Numbers: `3`, `3.14159`
2. Basic operations: `3 + 4`, `5 * 7`
3. Function application: `(f e1 ... en)`. Parentheses can be omitted if the context allows.
4. Conditional expressions: `(if b then e1 else e2)`

Haskell almost looks like a scripting language. In contrast to script languages like PHP, Ruby or Python, Haskell has all elements to realize even large software systems. In particular, unlike scripting languages, Haskell is *strictly typed*, that is, all values and expressions have a type that is checked by the Haskell system *before* the program is

### 3 Functional Programming

executed. We will discuss the different data types in more detail in chapter 3.2. Here, we want to annotate the type for all functions to make clear what they are used for.<sup>1</sup>

A basic data type is `Int`, the set of integers (or a finite subset of it). The type of a function is annotated by “`::`”, where several argument types and the result type are separated by “`->`”. Furthermore, the intuitive meaning of functions should be explained by a *comment* before the function definition, where comments are introduced by two minus signs and reach until the end of the line.

As an example, consider a function for calculating squares. We can define this in Haskell as follows.

---

```
-- Computes the square of a number.
square :: Int -> Int
square x = x * x
```

---

If this definition is stored in the file `Square.hs`, then you can use the Haskell interpreter `ghci`, which belongs to the Glasgow Haskell Compiler (GHC), as follows to use this function.

---

```
> ghci
GHCi, version 7.10.3: http://www.haskell.org/ghc/ :? for help
Prelude> :l Square
[1 of 1] Compiling Main (Square.hs, interpreted)
Ok, modules loaded: Main.
*Main> square 3
9
*Main> square (3 + 1)
16
*Main> :q
Leaving GHCi.
```

---

A function for calculating the minimum of two numbers can be defined in Haskell in this way.

---

```
-- Computes the minimum of two numbers.
min :: Int -> Int -> Int
min x y = if x <= y then x else y
```

---

---

<sup>1</sup>Haskell can, unlike many other languages, infer function types, so that it is not necessary write the type down. Nevertheless, it is a better programming style to add function types for program documentation.

### 3 Functional Programming

Next, we have look at the factorial function. It is defined mathematically as follows.

$$n! = \begin{cases} 1, & \text{falls } n = 0 \\ n \cdot (n - 1)!, & \text{otherwise} \end{cases}$$

The Haskell implementation looks like this.

---

```
-- Computes the factorial of a non-negative number.
fac :: Int -> Int
fac n = if n == 0 then 1 else n * fac (n - 1)
```

---

#### 3.1.1 Evaluation

The evaluation of function definitions in Haskell is done by oriented calculation from left to right: First the current parameters are bound, that is, the formal parameters are replaced by the current ones. Then the left side is replaced by the right side.




Figure 3.1: Ways to evaluate a function

Figure 3.1 shows how functions can be evaluated. The right branch shows how a function is evaluated in Java, the left branch resembles the evaluation in Haskell if double calculations of  $3 + 1$  are omitted.

### 3 Functional Programming

Another example is the evaluation of a call to our function `fac`.

$$\text{fac } 2 = \text{if } 2 == 0 \text{ then } 1 \text{ else } 2 * \text{fac } (2 - 1) \quad (3.1)$$

$$= \text{if False then } 1 \text{ else } 2 * \text{fac } (2 - 1) \quad (3.2)$$

$$= 2 * \text{fac } (2 - 1) \quad (3.3)$$

$$= 2 * \text{fac } 1 \quad (3.4)$$

$$= 2 * (\text{if } 1 == 0 \text{ then } 1 \text{ else } 1 * \text{fac } (1 - 1)) \quad (3.5)$$

$$= 2 * (\text{if False then } 1 \text{ else } 1 * \text{fac } (1 - 1)) \quad (3.6)$$

$$= 2 * 1 * \text{fac } (1 - 1) \quad (3.7)$$

$$= 2 * 1 * \text{fac } 0 \quad (3.8)$$

$$= 2 * 1 * (\text{if } 0 == 0 \text{ then } 1 \text{ else } 0 * \text{fac } (0 - 1)) \quad (3.9)$$

$$= 2 * 1 * (\text{if True then } 1 \text{ else } 0 * \text{fac } (0 - 1)) \quad (3.10)$$

$$= 2 * 1 * 1 \quad (3.11)$$

$$= 2 * 1 \quad (3.12)$$

$$= 2 \quad (3.13)$$

We now want to develop an efficient function for calculating Fibonacci numbers. Our first version is directly motivated by the mathematical definition.

---

```
-- Computes the n-th Fibonacci number.
fib1 :: Int -> Int
fib1 n = if n == 0
 then 0
 else if n == 1
 then 1
 else fib1 (n - 1) + fib1 (n - 2)
```

---

However, this variant is extremely inefficient: its execution time is  $O(2^n)$ .

How can we improve our first version? We calculate the Fibonacci numbers from the bottom: The numbers are enumerated from 0 until the  $n$ th number is reached: 0 1 1 2 3 ...  $\text{fib}(n)$ .

This programming technique is known by the name of *accumulator technique*. To do this, we must always keep the two previous numbers as parameters.

---

```
-- An accumulator function to compute n-th Fibonacci number more efficiently.
fib2' :: Int -> Int -> Int -> Int
fib2' fibn fibnp1 n = if n == 0
 then fibn
 else fib2' fibnp1 (fibn + fibnp1) (n - 1)
```

---

### 3 Functional Programming

```
-- Computes the n-th Fibonacci number.
fib2 :: Int -> Int
fib2 n = fib2' 0 1 n
```

---

Here, `fibn` is the  $n$ th Fibonacci number and `fibnp1` the  $(n + 1)$ th. Thus we achieve a linear runtime.

From a software-technical point of view, our second variant is unattractive: We want to avoid other external calls to `fib2'`. How this works is described in the next section.

#### 3.1.2 Local definitions

Haskell offers several ways to define functions locally. One possibility is the keyword `where`:

---

```
fib2 :: Int -> Int
fib2 n = fib2' 0 1 n
 where fib2' fibn fibnp1 n =
 if n == 0
 then fibn
 else fib2' fibnp1 (fibn + fibnp1) (n - 1)
```

---

`where` definitions are visible in the previous equation, outside they are invisible. Alternatively, the keyword `let` can be used.

---

```
fib2 n =
 let fib2' fibn fibnp1 n =
 if n == 0
 then fibn
 else fib2' fibnp1 (fibn + fibnp1) (n - 1)
 in fib2' 0 1 n
```

---

In contrast to `where`, `let` is an *expression*. The `fib2'` defined by `let` is only visible within the `let` expression.

`let ... in ...` can occur as an arbitrary expression.

---

```
(let x = 3
 y = 1
 in x + y) + 2
```

---

The above expression evaluates to 6.

The syntax of Haskell does not require parentheses and no separation of the individual definitions (for example by a semicolon) when defining such blocks. In Haskell, the *off-side rule* applies: The next symbol after `where` or `let` that is not a whitespace defines a




Figure 3.2: Off-side rule in Haskell

*block.*

- If the next line starts to the right of the block, it belongs to the same definition.
- If the next line starts at the edge of the block, a new definition in the block starts here.
- However, if the next line starts to the left of the block, the block before it ends here.

Local definitions offer a number of advantages.

- Name conflicts can be prevented
- Wrong usage of helper functions can be prevented
- Better readability
- Redundant computations can be avoided
- Helper functions with less parameters

Let us look at an example to avoid multiple calculations. Instead of the confusing function definition

---

```
f x y = y * (1 - y) + (1 + x * y) * (1 - y) + x * y
```

---

we write instead

---

```
f x y = let a = 1 - y
 b = x * y
 in y * a + (1 + b) * a + b
```

---

By using the local declaration constructs `let` and `where` it is also possible to save parameters for helper functions. This is illustrated by the following example.



### 3 Functional Programming

The predicate `isPrim` is supposed to check if the given number is a prime number. In Haskell we can express this as follows.

---

```
-- Checks whether a number is prime.
isPrim :: Int -> Bool
isPrim n = n /= 1 && checkDiv (div n 2)
 where checkDiv m =
 m == 1 || mod n m /= 0 && checkDiv (m - 1)
```

---

Here `/=` expresses inequality, `&&` stands for a conjunction, `||` for the logical or and `div` for integer division.

In the last line we do not need further parentheses, because `&&` binds stronger than `||`. The `n` is visible in `checkDiv`, because the latter is locally defined.

## 3.2 Data types

### 3.2.1 Basic data types

#### Integers

We have already used a basic data type of Haskell before: integers. In fact, Haskell distinguishes between two types of integers:

**Int:** Values  $-2^{31} \dots 2^{31} - 1$

**Integer:** arbitrary size (limited by memory)

Arithmetic operators: `+` `-` `*` `div` `mod`

Comparison operators: `<` `>` `<=` `>=` `==` `/=`

#### Booleans

Booleans are another basic data type.

**Bool:** `True` `False`

Operators: `&&` `||` `not` `==` `/=`

`==` means equivalence, `/=` exclusive or (XOR).

#### Floats

Floating point numbers are a basic data type as well.

Float: 0.3 -1.5e-2

Operators: similar to `Int`, but `/` instead of `div`, no `mod`

#### Chars

Unicode characters are a basic data type, too.

Char: `'a'` `'\n'` `'\NUL'` `'\214'`

The operators are defined in the library `Data.Char` and can be imported into the program by adding `import Data.Char` to the beginning of the Haskell program.

---

```
chr :: Int -> Char
ord :: Char -> Int
```

---

The operator `“::”` is used for optional type annotations of expressions, as explained below.

#### 3.2.2 Type annotations

As already mentioned, Haskell is a strictly typed programming language, that is, all values and expressions in Haskell have a type which can also be annotated.

---

```
3 :: Int
```

---

In this example, `3` is a value or expression, and `Int` is a type expression. Other examples of type annotations are as follows.

---

```
3 :: Integer
(3 == 4) || True :: Bool
(3 == (4 :: Int)) || True :: Bool
```

---

We can also specify type annotations for functions. These are written in a separate line:

---

```
square :: Int -> Int
square x = x * x
```

---

But what is the type of `min` (see section ??), considering that the function has two arguments?

---

```
min :: Int -> Int -> Int
```

---

A function arrow is therefore also used between the argument types written. We will see later why this is necessary.

#### 3.2.3 Algebraic data structures

Own data structures can be defined as new data types. Values are built using *constructors*. Constructors are freely interpreted functions and therefore cannot be reduced.

##### Defining an algebraic data type

Algebraic data types are defined in Haskell as follows.

---

```
data τ = $C_1 \tau_{11} \dots \tau_{1n_1}$ | ...
 | $C_k \tau_{k1} \dots \tau_{kn_k}$
```

---

where

- $\tau$  is the newly defined type,
- $C_1, \dots, C_k$  are the constructors and
- $\tau_{i1}, \dots, \tau_{in_i}$  are the argument types of the constructor  $C_i$ . Therefore

---


$$C_i :: \tau_{i1} \rightarrow \dots \rightarrow \tau_{in_i} \rightarrow \tau$$


---

holds.

Note: Both type and constructor names must begin with a capital letter in Haskell!

##### Examples

1. Enumerated types (only nullary constructors):

---

```
data Color = Red | Blue | Yellow
```

---

This data type defines three values of the type `Color`. `Bool` is also an enumerated type, predefined as

---

```
data Bool = False | True
```

---

2. Types with only one constructor:

---

```
data Complex = Complex Float Float
Complex 3.0 4.0 :: Complex
```

---

### 3 Functional Programming

---

This is allowed because Haskell works with separate namespaces for types and constructors.

How do we select individual components? In Haskell we use *pattern matching* instead of explicit selection functions (this possibility will be explained later):

---

```
-- Add two complex numbers:
addC :: Complex -> Complex -> Complex
addC (Complex r1 i1) (Complex r2 i2) = Complex (r1+r2) (i1+i2)
```

---

3. Lists (mixed types): For now, we consider only lists with elements of type `Int`:

---

```
data List = Nil | Cons Int List
```

---

`Nil` represents the empty list. The function `append` allows concatenating two lists.

---

```
-- Concatenate two lists of integer elements:
append :: List -> List -> List
append Nil ys = ys
append (Cons x xs) ys = Cons x (append xs ys)
```

---

This is a definition using several equations, where the first suitable one is selected. For example, a function call could be reduced in the following way.

---

```
append (Cons 1 (Cons 2 Nil)) (Cons 3 Nil)
= Cons 1 (append (Cons 2 Nil) (Cons 3 Nil))
= Cons 1 (Cons 2 (append Nil (Cons 3 Nil)))
= Cons 1 (Cons 2 (Cons 3 Nil))
```

---

Haskell's predefined lists look like this.

---

```
data [Int] = [] | Int:[Int]
```

---

`[]` corresponds to `Nil`, `“:”` is equivalent to `Cons` and `[Int]` is `List`.

The operator `“:”` is right-associative. Therefore, the following holds.

---

```
1:(2:(3:[]))
```

---

is equal to

---

```
1:2:3:[]
```

---

---

which can be written in the more compact way

---

```
[1,2,3]
```

---

Haskell also offers the operator `++` instead of `append`, predefined in the `Prelude`.

---

```
[] ++ ys = ys
(x:xs) ++ ys = x : xs ++ ys
```

---

*Operators* are binary functions that are used infix and begin with a special character. By means of parentheses, operators can be used like normal functions.

```
(++) :: [Int] -> [Int] -> [Int]
```

`[1] ++ [2]` is equal to `(++) [1] [2]`.

The other way around, binary functions can be used infix with single inverted commas `'...'`.

---

```
div 4 2
```

---

can be written as

---

```
4 `div` 2
```

---

For user-defined data types, it is not automatically possible to compare or output them. For this the keyword `deriving` can be added after the data type definition.

---

```
data MyType = ... deriving (Eq, Show, Ord)
```

---

## 3.3 Polymorphism

For the definition of universally usable data structures and operations Haskell supports *type polymorphism*. To explain this in more detail, we will determine the length of a list as an example:

---

```
-- Compute the length of a list of integers:
length :: [Int] -> Int
length [] = 0
```

### 3 Functional Programming

```
length (_,xs) = 1 + length xs
```

---

Of course, this definition also works for other lists, for example of type `[Char]`, `[Bool]` or `[[Int]]`.

Generally, one could say

---

```
length :: ∀ Type τ. [τ] -> Int
```

---

which is expressed in Haskell by means of type variables.

---

```
length :: [a] -> Int
```

---

What is the type of `(++)`?

---

```
(++) :: [a] -> [a] -> [a]
```

---

In consequence only lists with the same argument type can be concatenated. We consider another example.

---

```
-- Compute the last element of a list:
```

```
last :: [a] -> a
last [x] = x
last (x:xs) = last xs
```

---

This works because `[a]` is a list type, and `[x]` is a one-element list (corresponds to `x:[]`). However, we must not swap the two rules, otherwise every call to the function will yield a run-time error!

How can we define polymorphic data types ourselves? Haskell provides *type constructors* for the construction of types.

---

```
data K a1 ... am =
 C1 τ11 ... τ1n1
 | ...
 | Ck τk1 ... τknk
```

---

These data type definitions therefore look similar to the previous ones but here

- $K$  is a type constructor (not a type),
- $a_1, \dots, a_m$  are type variables and
- $\tau_{ik}$  are type expressions like basic types, type variables or a type constructor applied to type expressions.

### 3 Functional Programming

Functions and constructors are applied to values or expressions and generate values. Similarly, type constructors are applied to types and generate types.

As an example for a polymorphic datatype we want to model partial values in Haskell.

---

```
data Maybe a = Nothing | Just a
```

---

Then “`Maybe Int`” and also “`Maybe (Maybe Int)`” is a valid type.

If a type constructor in a type is applied to type variables, then the resulting type is also called *polymorphic*, as in the following example.

---

```
isNothing :: Maybe a -> Bool
isNothing Nothing = True
isNothing (Just _) = False
```

---

Another good example of a polymorphic datatype is the binary tree.

---

```
data Tree a = Leaf a | Node (Tree a) (Tree a)

-- Compute the height of a binary tree:
height :: Tree a -> Int
height (Leaf _) = 1
height (Node tl tr) = 1 + max (height tl) (height tr)
```

---

In Haskell, polymorphic lists are also predefined as syntactic sugar.

---

```
data [a] = [] | a : [a]
```

---

Although this definition is not a syntactically valid Haskell program, it can be understood as follows: The square brackets around `[a]` are the type constructor for lists, `a` is the element type, `[]` is the empty list and `:` is the list constructor.

For this reason, the following expressions are all the same.

---

```
(:) 1 ((:) 2 ((:) 3 []))
1 : (2 : (3 : []))
1 : 2 : 3 : []
[1,2,3]
```

---

According to the above definition, any complex list types such as `[Int]`, `[Maybe Bool]` or `[[Int]]` are possible.

Some functions on lists are for example `head`, `tail`, `last`, `concat` and `(!!)`.

---

```
head :: [a] -> a
head (x:_) = x

tail :: [a] -> [a]
tail (_:xs) = xs

last :: [a] -> a
last [x] = x
last (_:xs) = last xs

concat :: [[a]] -> [a]
concat [] = []
concat (l:ls) = l ++ concat ls

(!!) :: [a] -> Int -> a
(x:xs) !! n = if n == 0 then x
 else xs !! (n - 1)
```

---

The last function can also be defined as follows.

---

```
(x:_) !! 0 = x
(_:xs) !! n = xs !! (n - 1)
```

---

Strings are defined in Haskell as lists of characters.

---

```
type String = [Char]
```

---

Here, “**type**” initiates the definition of a *Typsynonyms*, that is, a new name (**String**) for another type expression (**[Char]**).

Thus the string "Hello" corresponds to the list 'H':'e': 'l': 'l': 'o': []. For this reason, all list functions also work for strings.

---

```
length ("Hello" ++ " folks!")
```

---

The above expression is evaluated 12.

Other predefined type constructors in Haskell are as follows.

- Sum of two types

---

```
data Either a b = Left a | Right b
```

---

This can be used, for example, to write values of “different types” to a list.



---

```
[Left 42, Right "Hallo"] :: [Either Int String]
```

---

Sum types can be processed via patterns, for example.

---

```
valOrLen :: Either Int String -> Int
valOrLen (Left v) = v
valOrLen (Right s) = length s
```

---

- Tuple

---

```
data (,) a b = (,) a b
data (,,) a b c = (,,) a b c
```

---

Some functions have already been defined for this as well.

---

```
(3,True) :: (Int,Bool)

fst :: (a, b) -> a
fst (x, _) = x

snd :: (a, b) -> b
snd (_, y) = y

zip :: [a] -> [b] -> [(a, b)]
zip [] _ = []
zip _ [] = []
zip (x:xs) (y:ys) = (x,y) : zip xs ys

unzip :: [(a, b)] -> ([a], [b])
unzip [] = ([], [])
unzip ((x,y):xys) = let (xs, ys) = unzip xys
 in (x:xs, y:ys)
```

---

In principle `unzip` is the inversion of `zip`, that is, if “`unzip zs`” evaluates to the result `(xs,ys)`, then “`zip xs ys`” evaluates again to `zs`. However, conversely `unzip (zip xs ys)` does not always evaluate to `(xs,ys)`!

## 3.4 Pattern Matching

As we have already seen, functions can be defined by several equations using *pattern matching*.

---

```
f pat11 ... pat1n = e1
⋮
f patk1 ... patkn = ek
```

---

This results in very concise programs, since you only have to define the right-hand sides for the special cases specified by the patterns. For multiple rules, Haskell selects and applies the textual first rule with a matching left side. Overlapping rules are allowed in principle, but they can lead to programs that are not very readable and should therefore be avoided.

#### 3.4.1 Structure of patterns

In principle, the patterns are data terms (that is, they contain no defined functions) with variables. The following patterns are possible.

- **x** (*Variable*): matches always, the variable is bound to the current value.
- **\_** (*Wildcard*): matches always, no binding.
- **C pat<sub>1</sub>...pat<sub>k</sub>** where **C** is a *k*-ary Constructor: matches, if the same constructor and arguments match with *pat<sub>1</sub>, ..., pat<sub>k</sub>*
- **x@pat** (*as pattern*): matches if **pat** matches; Additionally, **x** is bound to the whole matched term.

The *as pattern* also avoids overlapping patterns.

---

```
last :: [a] -> a
last [x] = x
last (x : xs@(_:_)) = last xs
```

---

In addition there are so-called *(n+k)* patterns for positive integers. We do not them explain here, because they are often not supported and not important.

Patterns can also be used with **let** and **where**:

---

```
unzip :: [(a, b)] -> ([a], [b])
unzip ((x,y) : xys) = (x:xs, y:ys)
 where
 (xs,ys) = unzip xys
```

---

#### 3.4.2 Case expressions

Sometimes it is also useful to branch expressions using pattern matching.

---

```

case e of pat1 -> e1
 :
 patn -> en

```

---

The above defines an expression with type  $e_1, \dots, e_n$ , which must all have the same type.  $e, pat_1, \dots, pat_n$  must also have the same type. After the keyword `of`, the off-side rule also applies, that is, the patterns  $pat_1, \dots, pat_n$  must all begin in the same column.

The result of `e` is matched against  $pat_1$  to  $pat_n$  one after the other. If a pattern matches, the whole `case` expression is replaced by the corresponding  $e_i$ .

As an example, consider extracting the lines of a string.

---

```

-- Breaks a string into a list of lines where a line is terminated at a
-- newline character. The resulting lines do not contain newline characters.
lines :: String -> [String]
lines "" = []
lines ('\n':cs) = "" : lines cs
lines (c:cs) = case lines cs of
 [] -> [[c]]
 (l:ls) -> (c : l) : ls

```

---

#### 3.4.3 Guards

Each pattern can have an additional boolean condition, also called *guard*.

---

```

fac :: Int -> Int
fac n | n == 0 = 1
 | otherwise = n * fac (n - 1)

```

---

`otherwise` is not a keyword, but a function that always evaluates to `True`.

By combining guards and case expressions, the first  $n$  elements of a list can be extracted.

---

```

-- Returns prefix of length n.
take :: Int -> [a] -> [a]
take n xs | n <= 0 = []
 | otherwise = case xs of
 [] -> []
 (x:xs) -> x : take (n-1) xs

```

---

Guards are also allowed for `let`, `where` and `case`.

## 3.5 Higher Order Functions

In Haskell, functions are not only used to define calculation methods, but functions are also "first class citizens" because they can be used like all other values (for example numbers). This means that functions can appear in data structures and also as parameters or results of other functions. In the latter case one speaks of "higher-order functions".. Higher-order functions can be used for

- generic programming
- defining program schemes (control structures)

This enables us to achieve a better reusability and a higher modularity of the program code.

### 3.5.1 Example: derivative function

The derivative function is a function that returns a new function for a function. The numerical calculation looks like this.

$$f'(x) = \lim_{dx \rightarrow 0} \frac{f(x+dx) - f(x)}{dx}$$

An implementation with a small  $dx$  could look like the following.

---

```
dx :: Float
dx = 0.0001

-- Computes the derivation of a (continuous) function.
derive :: (Float -> Float) -> (Float -> Float)
derive f = f'
where f' :: Float -> Float
 f' x = (f (x + dx) - f x) / dx
```

---

Now, `(derive sin) 0.0` evaluates to 1.0 and `(derive square) 1.0` evaluates to 2.00033.

### 3.5.2 Anonymous functions (lambda abstraction)

Sometimes one does not want to give every defined function a name (like `square`), but to write it directly where it is needed, as shown in the example below.

---

```
derive (\x -> x * x)
```

---

The argument corresponds to the function  $x \mapsto x^2$ . Such a function without name is also called *lambda abstraction* or *anonymous function*. Here `\` stands for  $\lambda$ , `x` is a parameter and `x * x` is an expression (the body of the function).

### 3 Functional Programming

In general, anonymous functions in Haskell are defined as follows.

---

```
\ p1 ... pn -> e
```

---

where  $p_1, \dots, p_n$  are patterns and  $e$  is an expression.

We can also write the following.

---

```
derive f = \x -> (f (x + dx) - f x) / dx
```

---

When using this function, we can observe that it behaves approximately like the derived function ( $y \mapsto 2y$ ).

---

```
(derive (\x -> x * x)) 0.0 -> 0.0001
(derive (\x -> x * x)) 2.0 -> 4.0001
(derive (\x -> x * x)) 4.0 -> 8.0001
```

---

But `derive` is not the only function with a functional result. For example, the function `add` can be defined in three different ways.

---

```
add :: Int -> Int -> Int
add x y = x + y
```

---

or

---

```
add = \x y -> x + y
```

---

or

---

```
add x = \y -> x + y
```

---

So `add` can also be seen as a constant that returns a function as a result, or as a function that takes an `Int` and returns a function that takes another `Int` and only then returns an `Int`.

Therefore, the types `Int → Int → Int` and `Int → (Int → Int)` must be identical. In fact, the type constructor “ $\rightarrow$ ” is defined as right-associative, so that this binding always applies if no brackets are written. You should note that  $(a \rightarrow b) \rightarrow c$  *not* the same as  $a \rightarrow b \rightarrow c$  or  $a \rightarrow (b \rightarrow c)$ !

So it would make sense to write the following independently of the definition of `add`.

---

```
derive (add 2)
```

---

If a function is applied to "too few" arguments, this is called *partial application*, partial. The partial application is made possible syntactically by currying. The name *currying* goes back to *Haskell B. Curry*, who discovered the following isomorphy in the 1940s:

$$[A \times B \rightarrow C] \simeq [A \rightarrow (B \rightarrow C)]$$

This means that a function with two arguments can also be interpreted as a function with one argument, which then returns another function for the second argument.

Actually, this technique was established much earlier by *Moses Schönfinkel* [16] in 1924. But because this article was published in German and "*Schönfinkeling*" cannot be pronounced so well in English, the term "Currying" has prevailed.

A number of functions can now be defined with the help of partial applications.

- `take 42 :: [a] -> [a]` yields the first up to 42 elements of a list.
- `(+) 1 :: Int -> Int` is the increment function.

For operators, so-called *sections* offer an additional, shortened notation.

- `(1+)` is the increment function.
- `(2-)` is equal to `\x -> 2-x`.
- `(/2)` is equal to `\x -> x/2`.
- `(-2)` is *not* equal to `\x -> x-2`, because the compiler cannot distinguish the minus sign from the unary minus operator.

Therefore, operators can also be partially applied to the second argument.

`(/b) a = (a/) b = a / b`

The order of arguments is a design decision because of partial application. The order can be modified with  $\lambda$ -abstractions and the function `flip`.

---

```
flip :: (a -> b -> c) -> b -> a -> c
flip f = \x y -> f y x
```

---

The function `flip` can be used, for example, with `take` in order to supply the list argument first.

---

```
(flip take) :: [a] -> Int -> [a]
(flip take) "Hello World!" :: Int -> [Char]
```

---

## 3.5.3 Generic Programming

We consider the following functions `incList` and `codeStr`.

---

```
incList :: [Int] -> [Int]
incList [] = []
incList (x:xs) = (x + 1) : incList xs

code :: Char -> Char
code c | c == 'Z' = 'A'
 | c == 'z' = 'a'
 | otherwise = chr (ord c + 1)

codeStr :: String -> String
codeStr "" = ""
codeStr (c:cs) = code c : codeStr cs
```

---

The expression `codeStr "Informatik"` evaluates to `"Jogpsnbujl"`. We observe that the definitions of `incList` and `codeStr` are nearly identical. Only the function that is applied to the list elements differs.

A generalized version is the function `map`.

---

```
map :: (a -> b) -> [a] -> [b]
map _ [] = []
map f (x:xs) = f x : map f xs
```

---

`incList` and `codeStr` can be defined more elegantly by using `map`.

---

```
incList = map (+1)
codeStr = map code
```

---

We look at another two examples: A function that yields the sum of all elements in a list and a function that calculates the sum of a string's Unicode values.

---

```
sum :: [Int] -> Int
sum [] = 0
sum (x:xs) = x + sum xs

checkSum :: String -> Int
checkSum "" = 1
checkSum (c:cs) = ord c + checkSum cs
```

---

Is there a common pattern? Indeed, both functions can be defined more easily by means

### 3 Functional Programming

of the powerful function `foldr`.

---

```
foldr :: (a -> b -> b) -> b -> [a] -> b
foldr _ e [] = e
foldr f e (x:xs) = f x (foldr f e xs)

sum = foldr (+) 0
checkSum = foldr (\c res -> ord c + res) 1
```

---

To understand `foldr`, the following perspective might help. The first argument of `foldr` of type `(a -> b -> b)` replaces the list constructor `(:)` in the list. The second argument of type `b` is the replacement for the empty list `[]`. Note that the type of the supplied function matches the type of `(:)`.

The following expressions are equivalent.

---

```
foldr f e [1,2,3]
= foldr f e ((:) 1 ((:) 2 ((:) 3 [])))
= (f 1 (f 2 (f 3 e)))
```

---

The general approach when designing such functions is as follows: Identify a common pattern and implement it with functional parameters.

Another pattern that we know is the function `filter` that filter elements with certain properties from a list.

---

```
filter :: (a -> Bool) -> [a] -> [a]
filter _ [] = []
filter p (x:xs) | p x = x : filter p xs
| otherwise = filter p xs
```

---

We can use `filter`, for example, to transform a list into a set, that is, removing all duplicates.:

---

```
nub :: [Int] -> [Int]
nub [] = []
nub (x:xs) = x : nub (filter (/= x) xs)
```

---

By using `filter`, we can implement *Quicksort* to sort lists.

---

```
qsort :: [Int] -> [Int]
qsort [] = []
qsort (x:xs) =
qsort (filter (<= x) xs) ++ [x] ++ qsort (filter (> x) xs)
```

---



---

`filter` can also be defined by using `foldr`.

---

```
filter p = foldr (\x ys -> if p x then x:ys else ys) []
```

---

The function `foldr` is a very generic skeleton, it is equivalent to a katamorphism in category theory.

Using `foldr` can have drawbacks sometimes. The expression

---

```
foldr (+) 0 [1,2,3] = 1 + (2 + (3 + 0))
```

---

leads to a large calculation that is first built up on the stack and then evaluated in the end.

An improved version can be implemented by means of the accumulator technique.

---

```
sum xs = sum' xs 0
where sum' :: [Int] -> Int -> Int
sum' [] s = s
sum' (x:xs) s = sum' xs (x + s)
```

---

The expression `sum [1,2,3]` is now being reduced to `((0 + 1) + 2) + 3`, which can be evaluated immediately.

An equivalent implementation can be achieved by using a different version of fold.

---

```
foldl :: (a -> b -> a) -> a -> [b] -> a
foldl _ e [] = e
foldl f e (x:xs) = foldl f (f e x) xs
```

---

Hence, a call `foldl f e (x1 : x2 : ... : xn : [])` is replaced with `f ... (f (f e x1) x2) ... xn`.

Now `sum` can be defined even simpler.

---

```
sum = foldl (+) 0
```

---

#### 3.5.4 Control structures

Many control structures that we know from other programming languages can be modeled in Haskell. First, we consider the `while` loop.

---

```
x = 1;
while x < 100
do
x = 2*x
od
```

---

In general, a `while` loop consists of the following.

- a state before executing the loop (initial value)
- a condition
- a loop body for changing the state

In Haskell, the `while` loop looks as follows.

---

```
while :: (a -> Bool) -> (a -> a) -> a -> a
while p f x | p x = while p f (f x)
| otherwise = x
```

---

For example, the expression `while (<100) (2*) 1` evaluates to 128.

Note that this is not a language extension! This control structure is nothing more than a function, a first class citizen.

#### 3.5.5 Functions as data

What are data structures? From an abstract point of view, data structures are objects with specific operations.

- Constructors (like `(:)` or `[]`)
- Selectors (like `head` or `tail`, and pattern matching)
- Test functions (like `null`, and pattern matching)
- Conjunctions (like `++`)

The important part is the functionality, that is, the interface, not the implementation. Therefore, a data structure corresponds to a set of functions.

As an example, we want to implement arrays with arbitrary elements in Haskell.

The constructs have the following type.

---

```
emptyArray :: Array a
putIndex :: Array a -> Int -> a -> Array a
```

---

Now we only need a single selector.

---

```
getIndex :: Array a -> Int -> a
```

---

---

Now we want to implement the interface. We can achieve this simply by not using other data structures, for example, lists or trees, but rather implement the array as a function. The implementation of this approach could look like the following.

---

```
type Array a = Int -> a

emptyArray i =
 error ("Access to non-initialized component " ++ show i)

getIndex a i = a i

putIndex a i v = a'
 where a' j | i == j = v
 | otherwise = a j
```

---

The advantage of this implementation is its conceptual clarity because the implementation is the same as the specification. One drawback is the access time that grows with an increasing number of `putIndex` calls.

#### 3.5.6 Useful Higher Order Functions

One useful higher-order function is the function composition operator `(.)`.

---

```
(.) :: (b -> c) -> (a -> b) -> a -> c
(f . g) x = f (g x)
```

---

Two more interesting higher-order functions are `curry` and `uncurry`. They allow using functions that are defined on tuples with single elements and vice versa.

---

```
curry :: ((a, b) -> c) -> a -> b -> c
curry f x y = f (x, y)

uncurry :: (a -> b -> c) -> (a, b) -> c
uncurry f (x, y) = f x y
```

---

Finally we have look at the function `const` that takes two arguments and returns the first one.

---

```
const :: a -> b -> a
const x _ = x
```

---

### 3.5.7 Higher Order Functions in Imperative Languages

Modern imperative programming languages allow functions as parameters or return values, too. Although partial application – like Haskell offers – is often not supported and the integration is not as seamless, many algorithms can be defined in a more compact way using higher-order functions. Higher-order also allows additional ways of abstraction since functions can not only abstract values but also program behavior.

In the following, we have a look at the usage of functional parameters in Ruby and Java 8. Besides the concrete syntax, we also discuss the problems that arise from the interplay of functions as values and mutable variables, objects or memory cells.

Higher-order functions are part of Ruby’s syntax. They are available in the form of “blocks”. A block is a sequence of commands that can be parameterized.

The simplest form of blocks are non-parameterized blocks like, for example, loop bodies. Such blocks can have parameters (noted between vertical lines), which makes them equivalent to anonymous functions. Methods and functions can, in addition to normal parameters, also have a block parameter. For some simple examples, we consider the methods `each` and `map` for arrays.

---

```
a = [1,2,3,4,5]

b = a.map do |x| x + 1 end

b.each do |x| puts x end
```

---

The program outputs the numbers from 2 to 6 on the screen by means of `puts`.<sup>2</sup>

The `map`-Method applies the supplied unary block to every element of the array. The original array is not modified but instead, an array of the same length is created. The `each` method also applies the supplied unary block to every element of the array but does not create a new one.

To understand how blocks are used, we want to define our own version of `map` that mutates the supplied array. In order to not get too deep into specific Ruby details at this time, we define the function not as a method but rather as an independent procedure `map!`<sup>3</sup> that uses the array as an additional parameter.

---

```
def map!(a)
 for i in 0..a.size-1 do
 a[i] = yield(a[i])
 end
end
```

---

<sup>2</sup>Blocks in Ruby can be noted alternatively by using curly brackets, for example `b = a.map {|x| x+1}`.

<sup>3</sup>The bang in the name of methods and functions is a Ruby custom that indicates that the method is mutating, that is, the supplied object is modified. Many methods exist in both versions.

### 3 Functional Programming

```
a = [1,2,3,4,5]

map!(a) do |x| x + 1 end

a.each do |x| puts x end
```

---

A block that is supplied to `map!` is applied with the keyword `yield`; in the above example it is applied to the argument (`a[i]`). The procedure `map!` is then applied to the array `a` and given the increment function block as an additional parameter.<sup>4</sup>

Thus, blocks are equivalent to functional parameters. Multiple blocks as parameters or returning a block is problematic. Therefore, Ruby allows transforming a block into a value of the class `Proc` by using `lambda`. The result can be used like any other value, that is, as an ordinary parameter or as part of a data structure.

Besides, the class `Proc` offers the method `call` that applies the function, or rather – the block, to parameters. As an example we define the function `foldr` with an explicit functional parameter instead of a block.

---

```
def foldr(f,e,a)
 if a == [] then
 e
 else
 f.call(a[0],foldr(f,e,a[1,a.size-1]))
 end
end
```

---

The sum of the elements of an array can then be calculated as follows.

---

```
puts foldr(lambda do |x,y| x+y end,0,[1,2,3,4,5,6,7,8,9,10])
```

---

Next, we want to define an array of functions. The *i*-th position of the array is supposed to be the function that adds the value *i* to its argument. Intuitively, this can be implemented like this.

---

```
a = [0,1,2,3,4,5,6,7,8,9]

for i in 0 .. a.size-1 do # iterate over array, begin with 0
 a[i] = lambda {|x| x+i } # Write respective increment function to i-th
end # position

puts a[3].call(70) # Apply function at index 3 to the value 70
```

---

<sup>4</sup>Ruby also offers a predefined method `map!`.

### 3 Functional Programming

When the program is executed, the result is unexpectedly not 73 but 79. This is because variables correspond to memory cells in Ruby which are modified while the program executes. The program does not add the value that `i` has in the loop body but instead when the function is executed, which is 9 in the example above. Thus, function bodies in imperative languages should not contain mutable variables. This can be achieved by, for example, creating the array of functions by using `map!`.

---

```
a = [0,1,2,3,4,5,6,7,8,9]

a.map! { |i| lambda {|x| x + i }}

puts a[3].call(70) # Hier erhalten wir nun 73.
```

---

Now the variable `i` is no longer a memory cell that can be modified over time. Instead, the variable is a block parameter that is created each time the block is applied, similar to scoping. The program now returns the expected value 73.

Anonymous functions can also be defined in Java, starting with version 8. Using functional parameters by means of anonymous inner classes was possible before but the new lambda notation increases the code readability drastically and adds to the feeling of functional programming. As an example we define the class `Higher` that offers a non-mutating `map` function for lists.

---

```
import java.util.*;

class Higher {

 interface Fun<A,B> { // define interface for objects of functional type
 B call (A arg); // that requires a call method
 }

 static <A,B> List map (Fun<A,B> f, List<A> xs) {
 List ys = new ArrayList (xs.size());
 for (A x : xs) {
 ys.add(f.call(x)); // the interface is used here
 }
 return ys;
 }
}
```

---

When using the `map` method, we now can use the lambda notation to define anonymous implementations of the `Fun` interface, as shown below.

---

```
public static void main (String[] args) {
```

### 3 Functional Programming

```
List<Integer> a = Arrays.asList(1,2,3,4,5);
a = map(x -> x + 1, a);
System.out.println(a); // Result: [2,3,4,5,6]
}
```

---

The lambda function `x -> x + 1`<sup>5</sup> is an elegant alternative to the definition of an anonymous inner class that implements the interface `Fun`.

We want to find out next, if the problem of using mutable variables within function bodies (as before in Ruby) is present in Java, too. Therefore, we again implement a loop that returns a list of increment functions.

---

```
public static void main (String[] args) {
 List<Fun<Integer,Integer>> fs = new ArrayList<>(10);
 for (int i = 0; i<10; i++) {
 fs.add(x -> x + i);
 }
 System.out.println(fs.get(3).call(70));
}
```

---

At compile time, the following error message appears.

---

```
local variables referenced from a lambda expression must be final or
effectively final
fs.add(x -> x + i);
```

---

Thus Java recognizes that the variable `i` is modified in the program and therefore is cannot be used in the body of a lambda expression. The simple solution is creating a `final` copy of `i` for every loop iteration.

---

```
public static void main (String[] args) {
 List<Fun<Integer,Integer>> fs = new ArrayList<Fun<Integer,Integer>>();
 for (int i = 0; i<10; i++) {
 final int j = i;
 fs.add(x -> x + j);
 }
 System.out.println(fs.get(3).call(70));
}
```

---

Now the program compiles and returns 73 as expected.

Modern imperative programming languages often offer the possibility to use functional

---

<sup>5</sup>Multiple function arguments are separated by commas within parentheses, for example `(x,y) -> x + y`. Type annotations like `int x -> x + 1` are also possible.

parameters and values, as demonstrated in the examples. The resulting code is often shorter and more comprehensible, especially when predefined functions like `map` and `fold` can be used for lists or arrays. The resulting way of programming often differs significantly from the familiar imperative way because control structures are less important and data plus function that manipulate data are moved to the foreground, that is, the core of the implemented algorithms.

We have also seen the major pitfall of using higher-order functions in imperative programming languages: mutable variables in function bodies. The problem can often be mitigated by not modifying variables in this context and programming in a more functional way. Especially in Java, using a `final` copy of a mutable variable is a solution, too.

## 3.6 Type Classes and Overloading

We consider the function `elem` that checks whether an element is contained in a list.

---

```
elem x [] = False
elem x (y:ys) = x == y || elem x ys
```

---

What are possible type of `elem`? Some examples are listed below.

---

```
Int -> [Int] -> Bool
Bool -> [Bool] -> Bool
Char -> String -> Bool
```

---

Unfortunately, “`a -> [a] -> Bool`” is not a valid type since an arbitrary type `a` is too general: `a` also includes functions, for which defining equality is difficult (there is no correct, general definition in Haskell). Thus, we need a way to restrict types to types for which value equality is defined. This can be written in Haskell as follows.

---

```
elem :: Eq a => a -> [a] -> Bool
```

---

“`Eq a`” is called a *type constraint*. By adding parentheses and commas, multiple type constraints can be used.

The class `Eq` is defined in Haskell like this.

---

```
class Eq a where
 (==), (/=) :: a -> a -> Bool
```

---

A *class* contains one or more functions that need to be defined for all instances of the



### 3 Functional Programming

class, that is, types. In case of `Eq`, the functions are `(==)` and `(/=)`. Types can be defined as *instances* of a class by implementing the functions of the class.

---

```
data IntTree = Empty | Node IntTree Int IntTree

instance Eq IntTree where
 Empty == Empty = True
 Node t1 n tr == Node t1' n' tr' = t1 == t1' && n == n'
 && tr == tr'
 - == - = False

t1 /= t2 = not (t1 == t2)
```

---

Now `(==)` and `(/=)` can be used for the type `IntTree` and, for example, the above function `elem` can also be used with lists of elements of the type `IntTree`.

Class instances can also be defined for polymorphic types. However, this might require type constraints when defining the instance, as shown in the next example.

---

```
data Tree a = Empty | Node (Tree a) a (Tree a)

instance Eq a => Eq (Tree a) where
 ...<as above>...
```

---

Note that infinitely many types become instances of the class `Eq` this way.

#### 3.6.1 Predefined Functions of a Class

The definition of `(/=)` looks like the above version for almost all instances. Therefore, a default definition is often useful for class functions. The default definition can be overwritten (and needs to be in some cases).

---

```
class Eq a where
 (==), (/=) :: a -> a -> Bool
 x1 == x2 = not (x1 /= x2)
 x1 /= x2 = not (x1 == x2)
```

---

#### 3.6.2 Predefined Classes

For some types it is useful to define a total order on values of the type. This functionality is provided in Haskell by an extension of `Eq`, .

---

```
data Ordering = LT | EQ | GT

class Eq a => Ord a where
 compare :: a -> a -> Ordering
 (<), (<=), (>=), (>) :: a -> a -> Bool
 max, min :: a -> a -> a

... -- predefined implementations
```

---

A minimal instance definition needs at least `compare` or `(<=)`.

Other predefined classes are `Num`, `Show` and `Read`.

- `Num` represents numbers for calculations. (`(+) :: Num a => a -> a -> a`)
- `Show` transforms values into strings. (`show :: Show a => a -> String`)
- `Read` constructs values from strings. (`read :: Read a => String -> a`)

Other predefined classes are presented in the lecture "Funktionale Programmierung" (functional programming).

Instances can be derived automatically (except for `Num`) for self-defined data types by appending the keyword `deriving` and a list of type classes to the data type definition.

---

```
deriving ($\kappa_1, ..., \kappa_n$)
```

---

*Exercise:* Check the types of all functions which were defined in the lecture for the most general type. This is possible by deleting the type signature for self-defined functions; Haskell will then derive the most general type. A few examples are the following.

---

```
(+) :: Num a => a -> a -> a
nub :: Eq a => [a] -> [a]
qsort :: Ord a => [a] -> [a]
```

---

#### 3.6.3 The class `Read`

The class `Show` is used to output data as text. To read data, that is, retrieve a value from a string, the string needs to be *parsed*, which is a difficult task. Luckily, there exists a predefined class. Nevertheless, using the class requires some understanding about parsing strings.

We consider the following type definition that defines the type of functions that parse strings to values.

---

```
type ReadS a = String -> [(a,String)]
```

---

---

What is the intention behind the return type of `ReadS`? The first part of the tuple is the actual result of the parser while the second part is the remainder of the string that has not been parsed yet. For example, when we consider the string `"Node Empty 42 Empty"`, it becomes clear that after reading the initial string `Node`, a tree needs to follow. In this situation, we would like to receive the remaining unparsed string to use it later.

If the string can be parsed multiple ways, the alternatives are returned as elements of a list. If the supplied string cannot be parsed, the returned value is the empty list.

Using `ReadS` can look like the following.

---

```
class Read a where
 readsPrec :: Int -> ReadS a
 readList :: ReadS [a] -- predefined
```

---

The two functions `reads` and `read` are defined as follows.

---

```
reads :: Read a => ReadS a
reads = readsPrec 0

read :: Read a => String -> a
read str = case reads str of
 [(x,"")] -> x
 _ -> error "no parse"
```

---

Thus, `reads` allows transforming a string into a value while checking whether the input is syntactically correct. On the other hand, `read` can be used when there is no doubt that the input is syntactically correct.

The evaluation of `reads` and `read` expressions looks like this.

---

```
reads "(3,'a')" :: [((Int,Char),String)]
= [((3,'a'),"")]

reads "(3,'a')" :: [((Int,Int),String)]
= []

read "(3,'a')" :: (Int,Char)
= (3,'a')

read "(3,'a')" :: (Int,Int)
= error: no parse

reads "3,'a'" :: [(Int,String)]
= [(3,",'a'")]
```

### 3.7 Evaluation Strategies and Lazy Evaluation

We consider the following Haskell program.

---

```
f x = 1
h = h
```

---

How is the expression `f h` being evaluated? Since the expression contains two function calls (`f` and `h`), it is unclear which one is evaluated first. If `h` is evaluated first, the program does not terminate, but if `f` is evaluated before `h`, the equation for `f` can be used to substitute the expression `f h` with `1`. Therefore, the order of evaluation is important! Although the results are always the same (because functional programming does not have side effects), whether or not the evaluation terminates for a certain result can differ.

Thus, which evaluation order should be used? The evaluation order is determined by the programming language, for example, Java would evaluate `h` first and thus never terminate but Haskell evaluates differently because the result is `1`. To precisely describe these differences, we define the term reduction strategy first and also give a formal definition of functional computation. Finally, we will discuss which strategy Haskell uses.

We already saw earlier that in a Haskell computations, expressions are being replaced with expressions that are considered equal according to the program rules. To define this process precisely, we need a few basic terms that are defined in the following. A few simplifications are made so that the definitions do not become too complex. In particular, type classes and polymorphism is not considered.

**Definition 3.1 (Program signature)** A program signature  $\Sigma = (S, F)$  consists of a set of sorts<sup>6</sup>  $S$  and a set of function signatures of the form<sup>7</sup>

$$f :: s_1 \dots s_n \rightarrow s$$

where  $f$  is a name,  $n \geq 0$  is a natural number and  $s_1, \dots, s_n, s \in S$ . In case of  $n = 0$  we sometimes omit the arrow, that is, we simply write  $f :: s$  instead of  $f :: \rightarrow s$ . Furthermore, for every name  $f$  in  $F$  only one such element exists.<sup>8</sup>

If  $f :: s_1 \dots s_n \rightarrow s \in F$  then  $f$  is called an  $n$ -ary function symbol.  $c$  is called a constant if  $c :: s \in F$ .

---

<sup>6</sup>Sorts are comparable to data types. The term "sort" is commonly used in mathematical logic.

<sup>7</sup>Since we also do not consider higher-order functions, we write function signatures in a non-curryfied way.

<sup>8</sup>This means that we do not allow overloading, as it would be possible with type classes.

### 3 Functional Programming

Thus, program signatures specify the types and function symbols that can be used in a program.  $\Sigma = (S, F)$  mit  $S = \{\text{Int}, \text{Bool}\}$  und

$$F = \{+ :: \text{Int Int} \rightarrow \text{Int}, * :: \text{Int Int} \rightarrow \text{Int}, \\ \text{square} :: \text{Int} \rightarrow \text{Int}, \text{True} :: \text{Bool}, \text{False} :: \text{Bool}, \dots\}$$

We assume that all literals are defined as constants in the program signature, for example is  $0 :: \rightarrow \text{Int} \in F$ .

Instead of expression, we also often use the word "term" because the vocabulary originates from the context of term substitution. Terms need to be structured logically, that is, they only contain variables and symbols from the program signature.

**Definition 3.2 (Term)** *Let  $\Sigma = (S, F)$  be a program signature and  $X$  a set of variables (of form  $x :: s$ ), which are disjunct from the symbols in  $\Sigma$ , that is, the variable names  $x$  are different from the function symbols. The set of terms  $T_\Sigma(X)_s$  of the signature  $s$  over  $\Sigma$  and  $X$  is the smallest set for which the following conditions hold.*

1. *For all constants  $c :: s \in F$  holds  $c \in T_\Sigma(X)_s$ .*
2. *For all variables  $x :: s \in X$  holds  $x \in T_\Sigma(X)_s$ .*
3. *For all  $n$ -ary functions  $f :: s_1 \dots s_n \rightarrow s \in F$  and  $t_i \in T_\Sigma(X)_{s_i}$ ,  $i = 1, \dots, n$ , is  $(f \ t_1 \dots t_n) \in T_\Sigma(X)_s$*

With  $T_\Sigma(X)$  we denote the set of all terms for a given signature, that is,

$$T_\Sigma(X) = \{t \mid s \in S, t \in T_\Sigma(X)_s\}$$

.

Note:

- This is an "inductive definition", that is, complex cases are defined based on simple cases. This is a common computer science approach to define complex structures.
- Intuitively, this definitions means that terms are constants, variables or a combination with appropriate arguments.
- Although, in the definition, the function symbol always comes first, we will use Haskell-like infix symbols in examples.

**Example:** Let  $\Sigma$  be as defined above,  $X = \{x :: \text{Int}, y :: \text{Int}, z :: \text{Int}\}$ . Then, the following holds.

$$(\text{square } x) \in T_\Sigma(X)_{\text{Int}}$$

$$(1 + 2) \in T_\Sigma(X)_{\text{Int}}$$

$$(\text{square } (\text{square } (y + z))) \in T_\Sigma(X)_{\text{Int}}$$

Thus, we now can precisely define what a (functional) program is.

**Definition 3.3 (Program, Term substitution system)** *A program or term substitution system is a triple  $(\Sigma, X, R)$  where*

### 3 Functional Programming

- $\Sigma$  is a program signature
- $X$  is set of variables that not overlap with  $\Sigma$
- $R$  is a set of rules of the form

$$f\ t_1 \dots t_n = r$$

where  $(f\ t_1 \dots t_n), r \in T_\Sigma(X)_s$  for a sort  $s$  (the result sort of the function  $f$ ). In  $r$  only variables that also occur in  $(f\ t_1 \dots t_n)$  may occur.  $f\ t_1 \dots t_n$  is called left-hand side and  $r$  right-hand side of the rule.

Therefore, a Haskell program can be seen as a term substitution system, although Haskell allows more than that.

When applying a rule, the variables of the left-hand side are substituted by other terms so that the rule matches. To describe this, we need to define the term substitution.

**Definition 3.4 (Substitution)** A substitution  $\sigma$  is a mapping  $\sigma : X \rightarrow T_\Sigma(X)$ , that maps every variable  $x :: s \in X$  to a term  $t \in T_\Sigma(X)_s$ .

The application of a substitution  $\sigma$  with a term  $t \in T_\Sigma(X)$ , written as  $t\sigma$ , is defined inductively as follows.

- If  $t :: s \in X$  then  $t\sigma = \sigma(t)$ .
- If  $t :: s \in \Sigma$  (that is,  $t$  is a constant), then  $t\sigma = t$ .
- If  $t = (f\ t_1 \dots t_n)$ , then  $t\sigma = (f\ t_1\sigma \dots t_n\sigma)$ .

Therefore, substitution only replaces variables – everything else remains untouched. Usually it is required that a substitution  $\sigma$  modifies only a finite number of variables, that is, the set

$$\{x \mid x :: s \in X \text{ mit } \sigma(x) \neq x\}$$

is finite. Then  $\sigma$  can be represented as the finite set of pairs of form

$$\{x \mapsto \sigma(x) \mid x :: s \in X \text{ mit } \sigma(x) \neq x\}$$

**Example:** Let  $\sigma = \{x \mapsto 2, y \mapsto (\text{square } z)\}$  be a substitution. Then the following holds.

$$(x * (y + z))\sigma = 2 * ((\text{square } z) + z)$$

To compute the value of a term, certain substitutions at certain positions within the term are performed. Therefore, we need to define how subterms of a term at a specific position are substituted (evaluated).

**Definition 3.5 (Position)** A position uniquely identifies a position within a term. If  $t$  is a term then  $\text{Pos}(t)$  is the set of all positions in  $t$ . Positions are commonly noted as sequences of natural numbers, where the empty sequence  $\langle \rangle$  denotes the root position, that is, the entry point of the term.

### 3 Functional Programming

**Example:** Let  $t = (* x (+ y z))$  and  $Pos(t) = \{\langle \rangle, \langle 1 \rangle, \langle 2 \rangle, \langle 2, 1 \rangle, \langle 2, 2 \rangle\}$ , then  $\langle \rangle$  is the entry point of the term.  $\langle 1 \rangle$  denotes the first argument, that is,  $x$ .  $\langle 2 \rangle$  denotes the subterm  $(+ y z)$  and  $\langle 2, 1 \rangle$  the  $y$  within the subterm. The  $z$  within the subterm has the position  $\langle 2, 2 \rangle$ .

**Definition 3.6 (Subterm)** If  $p \in Pos(t)$  then  $t|_p$  is called subterm of  $t$  at the position  $p$  and is defined as followed.

- If  $p = \langle \rangle$  then  $t|_{\langle \rangle} = t$
- If  $p = \langle p_1, \dots, p_k \rangle$  and  $t = (f t_1 \dots t_n)$  then  $t|_{\langle p_1, \dots, p_k \rangle} = t_{p_1}|_{\langle p_2, \dots, p_k \rangle}$

**Example:** Let  $t = (* x (+ y z))$  with  $t|_{\langle 2, 2 \rangle}$  to be determined.  $t|_{\langle 2 \rangle}$  would be  $(+ y z)$  and  $t|_{\langle 2, 2 \rangle}$  therefore only  $z$ .

**Definition 3.7 (Replacement)** If  $p \in Pos(t)$  and  $t'$  is a term, then  $t[t']_p$  is the replacement of  $t|_p$  with  $t'$  at the position  $p$  and is defined as follows.

- If  $p = \langle \rangle$  then  $t[t']_{\langle \rangle} = t'$
- If  $p = \langle p_1, \dots, p_k \rangle$  and  $t = (f t_1 \dots t_n)$  then

$$t[t']_p = (f t_1 \dots t_{p_1-1} t_{p_1}[t']_{\langle p_2 \dots p_k \rangle} t_{p_1+1} \dots t_n)$$

**Example:** Let  $t = (* x (+ y z))$  and  $t' = (+ z y)$ , then  $t[t']_{\langle 2 \rangle} = (* x (+ z y))$ .

Now we have all requirements for describing computations in the substitution model formally.

**Definition 3.8 (Reduction step)** Let  $P = (\Sigma, X, R)$  be a program. Then a reduction step  $t_1 \Rightarrow t_2$  is defined if  $p \in Pos(t_1)$ ,  $l = r \in R$  and a substitution  $\sigma$  exists with the property

- $l\sigma = t_1|_p$  ( $\sigma$  replaces formal parameters with the current ones)
- $t_2 = t_1[r\sigma]_p$  (replace subterms with right-hand side of the rule after substituting formal parameters)

A reduction or computation  $t_1 \Rightarrow^* t_2$  is a (possibly empty) sequence of steps

$$s_1 \Rightarrow s_2 \Rightarrow \dots \Rightarrow s_n$$

with  $s_1 = t_1$  and  $s_n = t_2$ .

**Example:**  $R$  contains the rule

---


$$(\text{square } x) = (* x x)$$


---

Then the following holds.

### 3 Functional Programming

---

`(square (+ 1 2))`  $\rightarrow$  `(* (+ 1 2) (+ 1 2))`

---

where  $p = \langle \rangle$  and  $\sigma(\mathbf{x}) = (+ 1 2)$ .

---

`(* 3 (square 4))`  $\rightarrow$  `(* 3 (* 4 4))`

---

where  $p = \langle 2 \rangle$  und  $\sigma(\mathbf{x}) = 4$ .

---

**Evaluation of predefined functions:** The evaluation of predefined functions like “+” or “\*” can be defined conceptionally by means of an infinite set of rules.

---

`(+ 0 0)`  $\rightarrow$  0  
`(+ 0 1)`  $\rightarrow$  1  
`(+ 0 2)`  $\rightarrow$  2  
 $\vdots$   
`(+ 1 2)`  $\rightarrow$  3  
 $\vdots$

---

Thus, `(square (+ 1 2))` can be evaluated the following way.

---

`(square (+ 1 2))`  
 $\rightarrow$  `(* (+ 1 2) (+ 1 2))`  
 $\rightarrow$  `(* 3 (+ 1 2))`  
 $\rightarrow$  `(* 3 3)`  
 $\rightarrow$  9

---

It is also possible to evaluate the term as follows.

---

`(square (+ 1 2))`  
 $\rightarrow$  `(square 3)`  
 $\rightarrow$  `(* 3 3)`  
 $\rightarrow$  9

---

We see that reductions yield different interim results but the final result is the same. Therefore, we first want to define what a final result is.

**Definition 3.9 (Normal form)** *A term  $t$  is in normal form if no other term  $t'$  exists with  $t \Rightarrow t'$ , that is, if no further reduction step is possible.*

Thus, normal forms correspond to final results or the values of expressions. The last example shows that different strategies exist to apply reductions. We consider two strategies in particular.



### Innermost/call-by-value/strict Evaluation

Informally, innermost means that the arguments of a function call are evaluated *before* the function is evaluated. The formal definition is as follows.

**Definition 3.10** *A reduction step  $t_1 \Rightarrow t_2$  with  $t_2 = t_1[r\sigma]_p$  is called innermost (call-by-value, strict evaluation or evaluation in applicative order) if  $t_1|_p = (f\ s_1 \dots s_n)$  and every  $s_i$  is in normal form.*

### Outermost/call-by-name/non-strict Evaluation

Informally, outermost means that the outer procedure calls are replaced before calls in the arguments until only calls to elemental procedures occur that can be evaluated. The formal definition is as follows.

**Definition 3.11** *A reduction step  $t_1 \Rightarrow t_2$  with  $t_2 = t_1[r\sigma]_p$  is called outermost (call-by-name, non-strict evaluation or evaluation in normal order) if  $t_1|_p$  is an outermost subterm where a reduction is possible. This means for  $p = \langle p_1 \dots p_k \rangle$  that at reduction at the positions  $\langle p_1 \dots p_i \rangle$  with  $i < k$  is not possible.*

In the above example, the second evaluation was strict (call-by-value) while the first one was non-strict (call-by-name).

---

```
(square (+ 1 2))
$ \rightarrow$ (* (+ 1 2) (+ 1 2))
$ \rightarrow$ (* 3 (+ 1 2))
$ \rightarrow$ (* 3 3)
$ \rightarrow$ 9
```

---

We can see that the evaluation of `(+ 1 2)` happens twice but that the final result is identical to the strict evaluation. However, non-strict evaluation can return a value for a term where strict evaluation does not terminate. The non-strict evaluation yields always a normal form if one exists. On the other hand, strict evaluation can often be implemented more efficiently. We revisit the first example.

---

```
f x = 1
h = h
```

---

The expression `f h` evaluates with non-strict evaluation to 1 while a strict evaluation does not terminate.

The expression `(+ (square 2) (square 5))` has multiple innermost positions. Thus, there are strategies that always pick, for example, the leftmost innermost function symbol.

### 3 Functional Programming

- *leftmost-innermost* (LI)
- *rightmost-innermost* (RI)
- *leftmost-outermost* (LO)
- *rightmost-outermost* (RO)

In general, a reduction strategy is a (partial) mapping of terms to positions that determines which positions are reduced next. Sometimes it can be useful to reduce multiple positions in parallel.

---

```
(+ (square 2) (square 5))
```

---

The above example can be reduced in parallel at the positions  $\langle 1 \rangle$  and  $\langle 2 \rangle$ . This is also called *parallel innermost* (PI) and *parallel outermost* (PO) reduction.

There are strict (LI) functional languages (Lisp, Scheme, Erlang, Standard ML) and non-strict (LO) functional languages (Haskell). In general, non-strict languages are *Berechnungsstärker*, as we saw before. However, the *Berechnungsstärke* can influence efficiency positively or negatively. The latter can happen when arguments are duplicated in a reduction step. Haskell avoids this by using lazy evaluation, as we will see later.

It is noteworthy that no completely strict programming language exists. At a minimum, the case distinction (if-then-else) needs to be evaluated non-strict. First, only the condition is evaluated and then, depending on the result, one of the alternatives is being evaluated.

One advantage of the outermost reduction is its *Berechnungsvollständigkeit*. Everything that can be computed in some way, will be computed by an PO strategy (usually also with a LO strategy but there are exceptions). In addition, a non-strict also has practical advantages because it supports

- avoiding unnecessary (possibly infinite) calculations and
- infinite data structures.

For example, the following function `from` defines the ascending, infinite list of natural numbers starting with `n`.

---

```
from :: Int -> [Int]
from n = n : from (n + 1)
```

---

We also consider the already familiar function `take`.

---

```
take :: Int -> [a] -> [a]
take n _ | n <= 0 = []
take _ [] = []
take n (x:xs) = x : take (n - 1) xs
```

---

`take 1 (from 1)` is evaluated to `[1]` because LO works as follows.

---

```
take 1 (from 1)
= take 1 (1:from 2)
= 1 : take 0 (from 2)
= 1 : []
```

---

The advantage lies within the separation of control (`take 1`) and data (`from 1`).

As another example, we have a look at calculating prime numbers via the sieve of Eratosthenes. The idea is as follows.

1. Consider the list of all numbers greater or equal to 2.
2. Remove all multiples of the first (prime) number.
3. The first element of the list is a prime number. Return to step 2. proceed the same with the remaining list.

This algorithm can be implement in Haskell as follows.

---

```
sieve :: [Int] -> [Int]
sieve (p:xs) = p : sieve (filter (\x -> x `mod` p > 0) xs)

primes :: [Int]
primes = sieve (from 2)
```

---

The argument of `sieve` is an input list which begins with a prime number and misses all multiples of smaller prime numbers. The result is the list of alle prime numbers!

Now the expression `take 10 primes` yields the first ten prime numbers.

---

```
[2,3,5,7,11,13,17,19,23,29]
```

---

By using `(!!)`, we can output the tenth prime number directly: `primes !! 9` evaluates to 29.

Infinite data structures can be used as an alternative to the accumulator technique. We consider the Fibonacci function as an example. To retrieve the  $n$ -th Fibonacci number, we create the list of all Fibonacci numbers and look up the  $n$ -th element.

---

```
fibgen :: Int -> Int -> [Int]
fibgen n1 n2 = n1 : fibgen n2 (n1 + n2)

fibs :: [Int]
fibs = fibgen 0 1

fib :: Int -> Int
fib n = fibs !! n
```

---

---

Now `fib 10` evaluates to 55.

Because infinite structures can often be useful, Haskell has some predefined functions to create infinite lists. For example, `repeat` yields an infinite list of identical elements.

---

```
repeat :: a -> [a]
repeat x = x : repeat x
```

---

Thus, `take 70 (repeat '-')` yields a textual line.

An infinite list of repeated application of a function can be created with the function `iterate`.

---

```
iterate :: (a -> a) -> a -> [a]
iterate f x = x : iterate f (f x)
```

---

As an exercise, one should think about what `iterate (+1) 0` evaluates to.

The example `from` has shown us that it is simple to create infinite arithmetic sequences in Haskell. We can also define *arithmetic sequences* and *intervals* with the step size 1 or arbitrary step size.

---

```
from :: Int -> [Int]
from n = n : from (n + 1)

fromThen :: Int -> Int -> [Int]
fromThen n1 n2 = let d = n2-n1 in n1 : fromThen (n1+d) (n2+d)

fromTo :: Int -> Int -> [Int]
fromTo n m = if n>m then [] else n : fromTo (n + 1) m

fromThenTo :: Int -> Int -> Int -> [Int]
fromThenTo n1 n2 m =
 let d = n2-n1
 in if d>=0 && n1>m || d<0 && n1<m
 then []
 else n1 : fromThenTo (n1+d) (n2+d) m
```

---

Since these functions<sup>9</sup> are rather useful, Haskell has a special syntax `[n..m]` for this functionality.

---

```
[n1 ..] $~\mbox{\textrm{means}}~$ from n
[n1,n2 ..] $~\mbox{\textrm{means}}~$ fromThen n1 n2
```

---

<sup>9</sup>In Haskell, these functions are defined with the prefix `enum`.

### 3 Functional Programming

```
[n .. m] $$$\mbox{\textrm{means}}$$$ fromTo n m
[n1,n2 .. m] $$$\mbox{\textrm{means}}$$$ fromThenTo n1 n2 m
```

---

The following equalities hold.

```
[1..4] == [1,2,3,4]
take 5 [2..] == [2,3,4,5,6]
take 5 [2,4..] == [2,4,6,8,10]
[1,3..10] == [1,3,5,7,9]
take 5 [3,1..] == [3,1,-1,-3,-5]
```

---

Not only integers but also floating point numbers and characters can be enumerated.

```
> [1,1.5 .. 10]
[1.0,1.5,2.0,2.5,3.0,3.5,4.0,4.5,5.0,5.5,6.0,6.5,7.0,7.5,8.0,8.5,9.0,9.5,10.0]
> ['a'..'z']
"abcdefghijklmnopqrstuvwxyz"
> take 20 ['A' ..]
"ABCDEFGHJKLMNOPQRST"
```

---

For this reason, Haskell's Prelude contains the type class `Enum`.

```
class Enum a where
 succ, pred :: a -> a
 toEnum :: Int -> a
 fromEnum :: a -> Int
 enumFrom :: a -> [a] -- [n..]
 enumFromThen :: a -> a -> [a] -- [n1,n2..]
 enumFromTo :: a -> a -> [a] -- [n..m]
 enumFromThenTo :: a -> a -> a -> [a] -- [n1,n2..m]
```

---

Therefore, the notation can be used for sequences of all instances of `Enum`. There are instances, for example, for `Int`, `Float`, `Char`, `Bool` and `Ordering`. Therefore, the following holds.

```
> [LT ..]
[LT,EQ,GT]
```

---

For data types with a finite number of members, the instance only yield finite sequences, too.

Arithmetic sequences are have many uses. For example, the factorial function can be defined as follows.

---

```
fac n = foldr (*) 1 [1 .. n]
```

---

It is also possible to enumerate a list of objects by pairing each element with an index.

---

```
> zip [1..] "abcde"
[(1,'a'),(2,'b'),(3,'c'),(4,'d'),(5,'e')]
```

---

The combination with other familiar functions is possible too, of course.

---

```
> map (uncurry (++)) (zip (map show [1..]) (take 5 (repeat ". Zeile")))
["1. Zeile","2. Zeile","3. Zeile","4. Zeile","5. Zeile"]
```

---

#### Sharing

One drawback of the LO strategy remains: Computations can be duplicated. We revisit the simple function `double`.

---

```
double x = x + x
```

---

If we supply `double 3` to the function, then the LI evaluation looks like this.

---

```
double (double 3) $ \Rightarrow$ double (3 + 3)
$ \Rightarrow$ double 6
$ \Rightarrow$ 6 + 6
$ \Rightarrow$ 12
```

---

The LO strategy yields this instead.

---

```
double (double 3) $ \Rightarrow$ double 3 + double 3
$ \Rightarrow$ (3 + 3) + double 3
$ \Rightarrow$ $$$ 6 + double 3
$ \Rightarrow$ $$$ 6 + (3 + 3)
$ \Rightarrow$ $$$ 6 + 6
$ \Rightarrow$ 12
```

---

Because of the obvious inefficiency, no real programming language uses LO.

One optimization of the strategy leads to *lazy evaluation*, where, instead of terms, graphs are being reduced. Variables of the program correspond to pointers to expressions and evaluating an expression updates the value of every variable that points to it. This is called *sharing*. Sharing can also be understood as normalization of the program, where

### 3 Functional Programming




Figure 3.3: Sharing bei lazy evaluation

for every subexpression a new variable is used. For the above example, this looks like this.

```
double x = x + x
```

```
main = let y = 3
 z = double y
 in double z
```

The evaluation works as shown in figure 3.3. Black lines indicate a reduction step and blue lines are pointers to expressions.

The programming language Haskell is based on lazy evaluation, that is, the combination




Figure 3.4: Zyklische Liste `ones`

of the LO strategy with sharing. This strategy was formalized by Launchbury in 1993 [10].

Lazy evaluation is optimal in respect to the length of the evaluation: There are no redundant computations as with LI and no duplicated expressions as with LO, although sometimes a lot of memory is needed.

Another advantage of lazy evaluation is that functions can be composed nicely: If we assume the existence of a producer function, for example, of the type `gen :: α → β`, and a consumer function, for example, with the type `con :: β → γ`. If we compose both functions `con . gen`, lazy evaluation does not create large interim data structures, but instead, only parts that are needed at the moment are stored in memory and freed as soon as they are no longer needed.

If, for example, we calculate the 100-th Fibonacci number by evaluating the expression `fibs !! 99`, only two numbers need to be kept in memory in principle. This also applies if we work with large files stepwise.

Haskell also allows *cyclic data structures* like the list

---

```
ones = 1 : ones
```

---

The list can be stored in a constant amount of memory, as shown in figure 3.4.



# Bibliography

- [1] S. Antoy and M. Hanus. Declarative programming with function patterns. In *Proceedings of the International Symposium on Logic-based Program Synthesis and Transformation (LOPSTR'05)*, pages 6–22. Springer LNCS 3901, 2005.
- [2] S. Antoy and M. Hanus. Functional logic programming. *Communications of the ACM*, 53(4):74–85, 2010.
- [3] T. Ball and B. Zorn. Teach foundational language principles. *Communications of the ACM*, 58(5):30–31, 2015.
- [4] L. Byrd. Understanding the control flow of Prolog programs. In *Proc. of the Workshop on Logic Programming*, Debrecen, 1980.
- [5] A. Church. A formulation of the simple theory of types. *Journal of Symbolic Logic*, 5:56–68, 1940.
- [6] K. Claessen and J. Hughes. QuickCheck: A lightweight tool for random testing of Haskell programs. In *International Conference on Functional Programming (ICFP'00)*, pages 268–279. ACM Press, 2000.
- [7] K.L. Clark. Negation as failure. In H. Gallaire and J. Minker, editors, *Logic and Data Bases*, pages 293–322. Plenum Press, 1978.
- [8] J. Corbin and M. Bidoit. A rehabilitation of Robinson’s unification algorithm. In *Proc. IFIP '83*, pages 909–914. North-Holland, 1983.
- [9] M. Hanus. Declarative processing of semistructured web data. In *Technical Communications of the 27th International Conference on Logic Programming*, volume 11, pages 198–208. Leibniz International Proceedings in Informatics (LIPIcs), 2011.
- [10] J. Launchbury. A natural semantics for lazy evaluation. In *Proc. 20th ACM Symposium on Principles of Programming Languages (POPL'93)*, pages 144–154. ACM Press, 1993.
- [11] A. Martelli and U. Montanari. An efficient unification algorithm. *ACM Transactions on Programming Languages and Systems*, 4(2):258–282, 1982.
- [12] M.S. Paterson and M.N. Wegman. Linear unification. *Journal of Computer and System Sciences*, 17:348–375, 1978.
- [13] S. Peyton Jones, editor. *Haskell 98 Language and Libraries—The Revised Report*. Cambridge University Press, 2003.
- [14] J.A. Robinson. A machine-oriented logic based on the resolution principle. *Journal of the ACM*, 12(1):23–41, 1965.
- [15] N. Savage. Using functions for easier programming. *Communications of the ACM*, 61(5):29–30, 2018.

## *Bibliography*

- [16] M. Schönfinkel. Über die Bausteine der mathematischen Logik. *Mathematische Annalen*, 92:305–316, 1924.

# List of Figures

|     |                                            |    |
|-----|--------------------------------------------|----|
| 1.1 | Expressions with wildcards . . . . .       | 5  |
| 2.1 | States of threads . . . . .                | 14 |
| 2.2 | Remote Method Invocation in Java . . . . . | 26 |
| 3.1 | Ways to evaluate a function . . . . .      | 33 |
| 3.2 | Off-side rule in Haskell . . . . .         | 36 |
| 3.3 | Sharing bei lazy evaluation . . . . .      | 75 |
| 3.4 | Zyklische Liste <b>ones</b> . . . . .      | 76 |